


How to use Selenium, successfully



The Selenium Guidebook

by Dave Haeffner

Preface

This book is not a full and comprehensive treatise that outlines every possible permutation of Selenium (the open-source software test automation tool for web applications). There are other books that already do this. My goal, instead, is to teach you the necessary skills to use Selenium successfully for your circumstance.

What you have before you is a distilled and actionable guide culled from my consulting practice and full time positions held doing automated software testing over the last six years.

My goal in writing this is to provide you with the materials I wish existed when I was starting out. I hope it serves you well.

What This Book Will Cover

This book focuses on the latest version of Selenium (a.k.a. Selenium WebDriver) and its use to test desktop web browsers. Selenium's predecessor, Selenium RC, will not be covered. Mobile testing with Appium will also not be covered. And while record-and-playback tools like Selenium IDE and Selenium Builder are great for getting started, they're abysmal for growing past that point. To that end, they will not be covered in this book either.

What will be covered? An approach on writing tests, in code, that are reliable, maintainable, and capable of be run on any browser/operating system combination and automated with continuous integration.

Who This Book Is For

This book is for anyone who wants to take automated acceptance testing seriously and isn't afraid to get their hands a little dirty. That is to say, this book is for anyone who wants to use computers for what they are good at, and free up you and people on your team to do what they are inherently good at (which does not include repetitive, mundane tasks). And don't worry if you're super new to programming, I'll cover the basics so you'll have a good platform to build from.

About The Examples In This Book

The examples in this book are written in Ruby. This is partly because that is the language in which I am strongest. But really, it has more to do with how readable and approachable the language is. By sticking with a scripting language that reads a lot like English, this book can be picked up by almost anyone regardless of their technical background and be immediately useful. And even if Ruby isn't your preferred language, the strategies and practices used are applicable regardless of your tech stack.

The application under test used in the examples throughout this book are from an open source

project of mine called the-internet. It's available [on GitHub](#) and is accessible [on Heroku](#).

The tests are written for version 3.4.0 of [RSpec](#) (a popular open source Behavior Driven Development Ruby testing framework).

All third-party libraries (a.k.a. "gems") are specified in a `Gemfile` and installed using [Bundler](#) with `bundle install`.

This book is accompanied by a zip file of all code examples referenced in the text. If you get stuck or want a quick reference on how something is supposed to look take a look there. There are a number of issues you may run into, and this is sure to help save you some time and headache.

How To Read This Book

Chapters 1 through 5 are focused on priming you with knowledge and questions to consider when it comes to test strategy, language selection, and test design. Chapter 6 is where we first start to code. From there, the examples build upon each other through chapter 16.

Chapter 17 is an accumulation of available resources that will help you find information on your own.

Code examples will correspond to each chapter number. For chapters with multiple parts, code examples will be broken up into numbered sub-folders (e.g., code examples referenced in Part 2 of Chapter 9 can be found in `code_examples/09_base_page_object/02/`).

Feedback

If you find a glitch in the book (e.g., grammar issue, code issue, etc.) or have questions/feedback please feel free to e-mail me at dhaeffner@gmail.com.

If you submit something and I end up using it in a future version of the book, I'll give you a shout-out in the acknowledgements.

Thanks and enjoy!

Table of Contents

1. [Selenium In A Nutshell](#)
2. [Defining A Test Strategy](#)
3. [Picking A Language](#)
4. [A Programming Primer](#)
5. [Anatomy Of A Good Acceptance Test](#)
6. [Writing Your First Test](#)
7. [Verifying Your Locators](#)
8. [Writing Re-usable Test Code](#)
9. [Writing Really Re-usable Test Code](#)
10. [Writing Resilient Test Code](#)
11. [Prepping For Use](#)
12. [Running A Different Browser Locally](#)
13. [Running Browsers In The Cloud](#)
14. [Speeding Up Your Test Runs](#)
15. [Simple Command-line Test Execution](#)
16. [Automating Your Test Runs](#)
17. [Finding Information On Your Own](#)
18. [Now You Are Ready](#)

Chapter 1

Selenium In A Nutshell

What Selenium Is and Is Not Good At

Selenium is a software robot sent from the future to help us test web applications. But keep in mind that it's not one of those fancy shape-shifting robots that can run really fast. Instead it's more like one of those really strong robots that's not very fast and is best suited for accomplishing a certain objective.

That is to say -- Selenium is really good at a specific set of things. If you know what those are and stick to them then you will be able to easily write reliable, scalable, and maintainable tests that you and your team can trust.

Selenium is built to automate browsers, specifically human interaction with them. Things like navigating to pages, clicking on things (e.g., buttons, links, etc.), typing into input fields, and dragging & dropping. It's less ideal for checking lower-level functionality like HTTP status codes or HTTP headers.

While you can use Selenium this way, it necessitates additional setup of a third-party tool (like [BrowserMob Proxy](#)), and it can be a bit of a rabbit hole to fall down since there are numerous test cases to consider at this level.

Selenium Highlights

Selenium works on every major browser, with a number of major programming languages, and on every major operating system. Each language binding and browser are actively being developed to stay current. Yes, even Internet Explorer (thanks to [Jim Evans](#) and [the Microsoft Edge team](#)!).

Selenium can be run on your local computer, on a server (with Selenium Remote), on your own set of servers (with Selenium Grid), or on a third-party cloud provider (like [Sauce Labs](#)). As your test suite grows, your test runs will take longer to complete. To speed them up you will need to run them in parallel, which is where the benefit of having your own servers or using a cloud provider comes. There's also the benefit of having numerous browser and operating system combinations at your disposal.

One of the guiding philosophies of Selenium is to be able to write your tests once and run them across multiple browsers. While this is a rosy proposition, it's not that simple. There are plenty of gotchas to watch out for. We'll cover these in detail throughout the book.

Chapter 2

Defining A Test Strategy

A great way to increase your chances of automated web testing success is to first map out a testing strategy. And the best way to do it is to answer four questions:

1. How does your business make money?
2. How do your users use your application?
3. What browsers are your users using?
4. What things have broken in the application before?

NOTE: For organizations that don't deal directly in dollars and cents (e.g., non-profits, federal agencies, etc.) you should focus on how the application generates value for the end user.

After answering these questions you should have an understanding of the functionality and browsers that matter to the application you're testing. This will help you narrow down your efforts to the things that matter most.

What To Do With The Answers

After answering these you should end up with a prioritized backlog of test automation tasks that you can work off of and track progress against.

Question 1: Money/Value

Every company's application makes money (or generates value) through core functionality that is defined by a series of increasingly-valuable user interactions -- a.k.a. a 'funnel'. Your answers to this question will help you determine what your funnel is.

This will be your highest priority for automation. Start an automation backlog to keep track.

Question 2: Usage Data

Odds are your application offers a robust set of functionality well beyond the funnel. Your answers to this question will help highlight what it is. And if you're basing these answers on usage data, then it will be broken down from highly used to lightly used. Tack these items onto your automation backlog below the answers from Question #1 (ordered by their frequency of use).

Question 3: Browsers

Now that you know what functionality is business critical and adopted by your users, you need to determine what browsers to focus your automated web testing efforts on. Your usage data will tell you this as well. It will help you determine which browsers you can reasonably avoid testing in

(e.g., based on non-existent or low usage numbers). Note the top 2 (or 3 depending on your numbers), but focus on the top 1 for now. This is the browser you will start using for automation.

Question 4: Risky Bits

To round out the strategy it is also good to think about what things have broken in the application before. To answer this question it's best to check your defect/bug tracker (if you have one) and to ask members of your team. What you end up with may read like a laundry list of browser specific issues or functionality that has been flaky or forgotten about.

Be sure to check the items in this list against the items in your automation backlog. If you find items that aren't listed on the automation backlog add them to the bottom. If you find items that are in the backlog, make a note in the backlog item that this has been an issue in the past.

For issues that have happened numerous times be sure to move the item up in backlog priority. And if find issues that keep cropping up on a specific browser, compare this browser to your short list from Question #3. If it's a browser that's not in your list but it's still important (e.g., a small pocket of valuable users), track it on the backlog, but put it at the bottom.

Now You Are Ready

This model works best for applications with existing functionality and does not speak directly to testing new functionality that's being actively developed. That's not to say that the two couldn't co-exist. It largely depends on your available resources and pace of development. But in order to reach high quality at speed, you first have to go slow.

By answering these questions you should now have a prioritized backlog of critical business functionality, a short list of browsers to focus on, and an understanding of the risky parts of your application to watch out for. With it you are on the right track because now you're focusing on the things that matter to the business and it's users.

Chapter 3

Picking A Language

In order to work well with Selenium, you need to choose a programming language to write your acceptance tests in. Conventional wisdom will tell you to choose the same language as what the application is written in. That way if you get stuck you can ask the developers on your team for help. But if you're not proficient in this language (or new to programming), then your progress will be slow and you'll likely end up asking for more developer help than they have time for -- hindering your automation efforts and setting you up for failure.

A great way to determine which language to go with is to answer one simple question: Who will own the automated tests?

Answering this question (and stepping through the discussion that unfolds from it) will help you more effectively choose the best programming language for your context.

What To Do With The Answer

If you're a tester (or a team of testers) and you will be building and maintaining the test automation, then it's best to consider what languages you (and your team) already have experience with (or are interested in gaining experience with) rather than what the application is built in. Your Selenium tests DO NOT have to be written in the same language as the application you are testing. Have a discussion with your team to see what interests them and go from there.

For example, there's not much advantage in writing your test automation in Java if your team has no previous development experience and will be the ones owning it. Instead, try to choose a language that is more approachable -- e.g., a scripting language like Ruby or Python.

But if you're a developer who is working on the application and just looking to add automated acceptance testing to the mix, then it makes sense to continue with the same language.

Some Additional Food For Thought

As your suite of tests grows you will find the need to add functionality to make things easier to use, run faster, offer better reporting, etc. When this happens you will effectively be doing two jobs, no longer just writing automated tests but also growing a framework (a.k.a. a test harness).

As you are considering which language to go with, consider what open source frameworks already exist for that language. Going with one will save you a lot of time and give you a host of functionality out of the box that you would otherwise have to build and maintain yourself -- and it's FREE.

You can see a list of open source Selenium WebDriver frameworks and their respective languages [here](#).

Outro

Choosing a programming language for automated testing is not a decision that should be taken lightly. If you're just starting out (or looking to port your tests) then considering and discussing these things will help position you for greater success.

With all that being said, the rest of this book will show you how to build your own test harness (in addition to writing well factored tests) in Ruby.

Chapter 4

A Programming Primer

This section will ply you with just enough programming concepts (and how they pertain to automated web testing) to get you going so that you have some working knowledge and a vocabulary that will enable you follow along with what you will see throughout this book and in your work after you put this book down.

Don't get too hung up on the details yet. If something doesn't make sense, it should once we dig into the full examples in the following chapters. Consider this more of a glossary than a code school.

Installation

Ruby is a quickly evolving ecosystem. In order to find the latest on how to install Ruby, I encourage you to [check out the official Ruby site's download page](#).

That being said, here are some install instructions to help you get started quickly.

- [Linux](#)
- [OSX](#)
- [Windows](#)

If you plan on doing serious development with Ruby, then I would consider using a version manager instead of a direct Ruby installation (e.g., [RVM](#) or [rbenv](#)). This will enable you to run different versions of Ruby at the same time on your machine as well as have different sets of dependencies. It's crucial if you plan to work with multiple projects in Ruby that have different dependencies. But it's not required for this book.

The examples in this book have been tested in Ruby 2.2.2p95. So as long as you have a fairly recent version (at least 2.2.x) you should be good to go.

Installing Third-Party Libraries

One of the main benefits of Ruby is that it has a vibrant open source community with copious libraries (a.k.a. "gems") immediately available, making it simple to build complex things quickly. You can find out more about gems on [the official Ruby site's libraries page](#).

To install a gem directly, you just have to type `gem install` and the gem name (from your command-line) -- e.g., `gem install selenium-webdriver`. If you get a permission error then you will need to start the command with `sudo` in order to temporarily elevate your terminal session to the correct level (e.g., `sudo gem install selenium-webdriver`). When this happens, you will be

prompted for your password. Once provided, the gem will install and you'll be able to move forward.

A helpful library called [Bundler](#) is the recommended way to manage gems within a project. This is what I use and what was used in building out the examples for this book. All you need to know is that if you see a `Gemfile` in a directory, just run `bundle install` from your terminal window (assuming you have already installed the bundler gem -- `gem install bundler`). When you do, all the necessary gems (and their dependencies) will be installed for you.

If after installing gems with Bundler you have trouble executing your test code, try prepending your execution with `bundle exec` (e.g., `bundle exec rspec`).

Interactive Prompt

One of the immediate advantages to using a scripting language like Ruby is that you get access to an interactive prompt. Just type `irb` (which stands for "interactive ruby shell") from the command-line. It will load a prompt that looks like this:

```
irb(main):001:0>
```

In this prompt you can type out Ruby code. It will evaluate it and return the result. As you step through this chapter it will serve as a great place to practice the commands and concepts you will see.

When you're done, just type `quit`.

Choosing A Text Editor

In order to write Ruby code, you will need to use a text editor. Some popular ones in the Ruby community are [Vim](#), [Emacs](#), [Sublime Text](#), and [RubyMine](#).

It's important to pick an editor that works for you and makes you productive. So if you're new to programming and text editors then it's probably best to go with something more intuitive like Sublime Text.

Programming Concepts In A Nutshell

Programming can be a deep and intimidating thing if you're new to it. But don't worry. When it comes to testing there is only a handful of concepts that we really need to get started. Granted, the more you know, the better off you'll be. But you don't need to know a whole lot to be an effective test automator right now.

Of all the programming concepts out there, here are the ones you'll need to pay attention to now:

- Object Structures (Variables, Methods, and Classes)

- Scope
- Object Types (Strings, Numbers, Collections, Booleans)
- Actions (Assertions, Conditionals, Iteration)
- Inheritance

Let's step through each and how they pertain to testing with Selenium.

Object Structures

Variables

Variables are places to store and retrieve values of various types (e.g., Strings, Integers, Booleans, Collections, etc). Variables are created and then referenced by their name.

A variable name:

- can be one or more words in length
- use an underbar (`_`) to separate the words (e.g., `example_variable`)
- start with a lowercase letter
- are often entirely lowercase

You can store things in them by using an equals sign (`=`) after their name. In Ruby, a variable takes on the type of the value you store in it (more on object types later).

```
example_variable = "42"
puts example_variable.class
# outputs: String

example_variable = 42
puts example_variable.class
# outputs: Fixnum
```

In the above example `puts` is used to output a message. This is a common command that is useful for generating output to the terminal.

In Selenium, a common example of a variable is when we need to store an element (or the value from an element) in the beginning of a test to be referenced later. For example, getting a page's title.

```
page_title = @driver.title
```

`@driver` is the variable we will use to interact with Selenium. More on why it uses `@` soon (in Scope).

Methods

One way to group common actions (a.k.a. behavior) for easy reuse is to place them into methods. We define a method with the opening keyword `def`, a name (in the same fashion as a variable), and close it with the keyword `end`. Referencing a method is done the same way as a variable -- by it's name.

```
def example_method
  # your code
  # goes here
end

example_method
```

Additionally, we can specify arguments we want to pass into the method when calling it.

```
def say(message)
  puts message
end

say 'Hello World!'

# outputs:
# Hello World!
```

When setting an argument, we can also set a default value to use if no argument is provided.

```
def say(message = 'Hello World!')
  puts message
end

say
say 'something else'

# outputs:
# Hello World!
# something else
```

We'll see this tactic used in Selenium when we are telling Selenium how to wait with explicit waits (more on that in Chapter 10).

```
def wait_for(seconds=8)
  Selenium::WebDriver::Wait.new(:timeout => seconds).until { yield }
end
```

Classes

Classes are a useful way to represent concepts that will be reused numerous times in multiple places. They can contain variables and methods and are defined with the word `class` followed by the name you wish to give it.

Class names:

- must start with a capital letter
- should be CamelCase for multiple words (e.g., `class ExampleClass`)
- should be descriptive

You first have to define a class, and then create an instance of it (a.k.a. instantiation) in order to use it. Once you have an instance you can access the methods within it to trigger an action.

```
class Message
  def say(message = 'Hello World!')
    puts message
  end
end

message_instance = Message.new
message_instance.say 'This is an instance of a class'

# outputs: This is an instance of a class
```

An example of this in Selenium is the representation of a web page -- also known as a 'Page Object'. In it you will store the page's elements and behavior.

```
class Login
  LOGIN_FORM      = { id: 'login' }
  USERNAME_INPUT  = { id: 'username' }
  PASSWORD_INPUT  = { id: 'password' }

  def with(username, password)
    ...
  end
end
```

The variables that are fully capitalized are called constants, the values in curly brackets (`{ }`) are called hashes, and they are using something called symbols. More on all of that soon.

Scope

Now that we have our different object structures it's worth briefly covering how they can and cannot access one another. The easiest way to do this is by discussing the different types of variables you are likely to run into.

Local Variables

Local variables enable you to store and retrieve values from within a confined area (this is no different than what was described in the variables section earlier). A variable that was declared within a method is a great example of this. It is useful within the method it was declared, but inaccessible outside of it.

In your Selenium tests, a local variable will only be available from within the test that it was created.

Instance Variables

Instance variables enable you to store and retrieve values more broadly (e.g., both inside and outside of methods). They are annotated the same way as regular variables, except that they start with `@`.

A common example you will see throughout this book is the usage of `@driver`. This is an instance of Selenium stored in an instance variable. This object is what enables us to control the browser and by storing it as an instance variable our tests can easily use it.

Constants

Constants are for storing information that will not change. They are easy to spot since they start with a capital letter, and are often all uppercase. They share similarities to instance variables since they can be accessed more broadly.

They are commonly used to store element locator information at the top of Page Objects.

```
class Login

  LOGIN_FORM      = { id: 'login' }
  USERNAME_INPUT  = { id: 'username' }
  PASSWORD_INPUT  = { id: 'password' }
  ...
```

Environment Variables

Environment variables are a way to pass information into Ruby from the command-line. They are also a way to make a value globally accessible (e.g., across an entire program). They can be set and retrieved from within your code by:

- starting with the keyword `ENV`
- specifying the name of the variable in brackets (`ENV[]`)
- surrounding the variable name with single-quotes (`ENV[' ']`)
- setting a value for the variable using an equals sign (`ENV['example_variable'] =`)

- specifying a string value to store in the variable (`ENV['example_variable'] = 'example value'`)

Environment variables are often used to store configuration values that could change. A great example of this is the base URL for the application you're testing.

```
ENV['base_url'] = 'http://the-internet.herokuapp.com'
```

To change the value when running your Ruby application, you just have to specify a new value before the application. `sh base_url='http://localhost:4567' rspec login_spec.rb`

Object Types

Strings

Strings are alpha-numeric characters packed together (e.g., letters, numbers, and most special characters) surrounded by either single (`' '`) or double (`" "`) quotes. Single quotes are encouraged unless you intend to manipulate the string value through a tactic called interpolation. Interpolation enables you to inject code into your string to create dynamic output.

```
motd = 'Hello World!'
puts "The message of the day is: #{motd}"

# outputs:
# The message of the day is: Hello World!
```

You'll also want to use double-quotes if you need to have a single quote in your string.

```
puts "How's this?"
```

You'll run into Strings when working with copy on a page (e.g., pulling a page's URL, title, or h1 tag to determine if your test is in the right place before proceeding).

Numbers

The two common types of numbers you will run into with testing are Fixnum (whole numbers or integers) and Float (decimals).

If you have a test that needs to pull some values from a page and add/subtract/multiply/divide them, then this will come in handy. Although you may need to convert the values from a String to an Fixnum first. But don't sweat it, this is a trivial thing to do in Ruby.

```
count = @driver.find_elements(id: '#count').text.to_i
```


The conversion from a String to an Fixnum is done with `.to_i` (the 'i' stands for Integer). If you're working with decimals, you can use `.to_f` to convert it to a Float instead.

Collections

Collections enable you to gather a set of data for later use. In Ruby there are two types of built-in collections -- Arrays and Hashes.

Array values are stored in an ordered list, each with an index number (which starts at 0), and are surrounded by square brackets (`[]`). Hashes store values in the order they were added and use a key/value pair to store and retrieve them. Hashes are surrounded by curly brackets (`{ }`).

Both Arrays and Hashes are able to receive values of any type.

```
# Array Example

an_array = ["one", 2, 3.0, "four"]
puts an_array[0].class
puts an_array[1].class
puts an_array[2].class
puts an_array[3].class

# outputs:
# String
# Fixnum
# Float
# String
```

The array has four elements and the count starts at `0` . So when we access each of the values in the array we start at `0` and go till `3` .

You'll end up working with Arrays (or something similar) if you need to test things like HTML data tables, drop-down lists, or if you need to take an action against a specific element within a large list but there are no specific locators for it.

```
# Hash Example

a_hash = {one: "one", two: 2, three: [3]}
puts a_hash[:one].class
puts a_hash[:two].class
puts a_hash[:three].class

# outputs:
String
Fixnum
Array
```

Note that in the hash we are storing a string, a fixnum, and an array. And we are accessing each of the values by their key. Which in this case starts with a colon (:). This type of object is known as a Symbol, and it is often used as an identifier for objects in Ruby.

You'll end up working with Hashes and Symbols in your Page Objects to store and retrieve your page's locators.

```
class Login

  LOGIN_FORM      = { id: 'login' }
  USERNAME_INPUT  = { id: 'username' }
  PASSWORD_INPUT  = { id: 'password' }
  ...
```

Booleans

Booleans are binary values that are returned when asking a question of your code. They are what enable us to complete assertions.

There are numerous ways to ask questions. Some involve various [comparison operators](#) (e.g., `==`, `!=`, `<`, `>`, `<=>`), and others end in a question mark (e.g., `include?`). The response is either `true` or `false`.

```
@driver.get 'http://the-internet.herokuapp.com'
@driver.title.include?('The Internet')

# returns: true
```

Actions

Assertions

Assertions are made against booleans and result in either a passing or failing test. In order to leverage assertions we will need to use a testing framework (e.g., [RSpec](#), [minitest](#), or [test-unit](#)). For the examples in this book we will be using RSpec (version 3.4.0).

RSpec enables easy to read assertions through its [built-in matchers](#). With them our assertions will start with the word `expect` and end with `.to` followed by things like `equal` (or `eq`), `include`, or `be`.

We use these matchers by calling them on the variable we want to ask a question of.

```
@driver.get 'http://the-internet.herokuapp.com'
expect(@driver.title).to include('The Internet')

# or

@driver.get 'http://the-internet.herokuapp.com'
title_present? = @driver.title.include?('The Internet')
expect(title_present?).to eq true
```

Both approaches will work, resulting in a passing assertion. If this is the only assertion in your test then this will result in a passing test. More on good test writing practices in Chapter 5.

Conditionals

Conditionals work with booleans as well. They enable you execute different code paths based on their values.

The most common conditionals in Ruby are `if` and `case` statements. They both accomplish the same thing. They are just stylistically different approaches. Which approach you end up going with initially is really just a matter of preference.

```
number = 10
if number > 10
  puts 'The number is greater than 10'
elsif number < 10
  puts 'The number is less than 10'
elsif number == 10
  puts 'The number is 10'
else
  puts "I don't know what the number is."
end

# outputs: The number is 10
```

Note that in order to do an else/if statement it is `elsif`, not `elseif`.

```

number = 10
case number
when 11..100
  puts 'The number is greater than 10'
when 0..9
  puts 'The number is less than 10'
when 10
  puts 'The number is 10'
else
  puts "I don't know what the number is."
end

# outputs: The number is 10

```

You can do greater-than (`>`) and less-than (`<`) comparisons in a `case` statement as well, but it ends up looking a lot like an `if` block which takes away from the simplicity of the `case` statement.

You'll end up using conditionals in your test setup code to determine which browser to load based on an environment variable. Or whether or not to run your tests locally or somewhere else.

```

config.before(:each) do
  case ENV['browser']
  when 'firefox'
    @driver = Selenium::WebDriver.for :firefox
  when 'chrome'
    @driver = Selenium::WebDriver.for :chrome
  end
end

```

Iteration

Collections wouldn't be nearly as valuable without the ability to iterate over them one at a time. And in Ruby, it's simple to do. The syntax may initially feel awkward, but once you get the hang of it it's pretty straightforward.

Each collection comes enabled with methods for iteration. The most common one is `.each`. When using it you will need to specify a variable and a block of code.

The variable will be enclosed in pipes (`| |`) and represent the item of the collection that is being referenced one at a time. The variable will only be accessible within the code block. The block of code will open and close with the keywords `do` and `end`. In between the keywords is where you will put the code you want to execute (which will interact with the variable).

```
numbers = [1,2,3,4,5]
numbers.each do |number|
  puts number
end

# outputs:
# 1
# 2
# 3
# 4
# 5
# 6
# 7
# 8
# 9
# 10
```

After setting numbers to an array we are able to iterate over it and display each of its values with `puts`.

We can take this approach and couple it with a conditional to influence the output.

```
numbers = [1,2,3,4,5,6,7,8,9,10]
numbers.each do |number|
  puts number if number > 5
end

# outputs:
# 6
# 7
# 8
# 9
# 10
```

Iteration will come in handy in your Selenium tests if you have to loop over a collection of page elements to interact with them (e.g., HTML data tables, drop-down lists, etc.).

Inheritance

Classes have the ability to connect to one-another through parent/child inheritance. By having a single parent class we can store common actions in methods that can be readily available to all child classes.

Inheritance is done when declaring a child class by:

- providing the class name
- using a less-than symbol (<)
- providing the name of the parent class

```
class Parent
  def hair_color
    puts 'Brown'
  end
end

class Child < Parent
end

child = Child.new
puts child.hair_color

# outputs: Brown
```

You'll see this in your tests when writing all of the common Selenium actions you intend to use into methods in a parent class. Inheriting this class will allow you to call these methods in your child classes (more on this in Chapter 9).

Additional Resources

If you want to dive deeper into Ruby then I encourage you to check out some of the following resources:

- [Ruby Monk](#)
- [Codecademy](#)
- [Learn To Program](#)
- [The Pick-axe book](#)
- [The Pragmatic Studio's Online Course](#)

Chapter 5

Anatomy Of A Good Acceptance Test

In order to write acceptance tests that perform well and are both maintainable and resilient there are some simple guidelines to follow:

- Write atomic and autonomous tests
- Group similar tests together in small batches
- Be descriptive
- Use a Test Framework
- Store tests in a Version Control System

Atomic & Autonomous Tests

Each test needs to be concise (e.g., testing a single feature rather than multiple features) and capable of being run independently (e.g., sets up its own data rather than relying on a previous test to do it). Doing this may require a mental shift, discipline, and more up front effort. But it will make a dramatic impact on the quality, effectiveness, and maintainability of your tests.

Grouping Tests

As your test suite grows you should have multiple test files, each containing a small grouping of tests broken out by functionality that they're exercising. This will go a long way towards organization and maintenance as your test suite grows -- as well as faster execution times (depending on your approach to parallelization).

Being Descriptive

Each test file should be named appropriately, and each test within it should have an informative name (even if it may be a bit verbose). Also, each test (or grouping of tests) should be tagged with some helpful information to provide context and enable more flexible test execution (more on Tagging in Chapter 16).

This way all or part of your test suite can be run, and the results will be informative thanks to helpful naming.

This approach is useful for both humans and robots alike. It enables developers to run a subset of tests to exercise functionality they just modified (as part of their pre-check-in testing) while also enabling you to wire your suite up to a Continuous Integration server for fast and dynamic feedback that's fully automated (more on this in Chapter 16).

Test Framework

At the heart of every test suite is some kind of a test framework that does a lot of the heavy lifting (e.g., assertions, test group execution, easy global configuration for setup and teardown, reporting, etc.). Rather than reinvent the wheel, you can use one of the many that already exists (there's more than one for every programming language). And with it you can bolt on third party libraries to extend its functionality if there's something missing.

The examples in this book use [RSpec](#), a popular open source Behavior Driven Development testing framework for Ruby.

Version Control

In order to effectively collaborate with other testers and developers, your test code must live in a version control system of some sort. Look to see what your development team uses and add your code to it. Otherwise, set up one of the following:

- [Git](#)
- [Mercurial](#)
- [Subversion](#)

Keep in mind that your test code can live in a separate repository from the code of the application you're testing. Combining them may be advantageous, but if all you're doing is writing and running tests against web endpoints (which is a majority of what your testing will be with Selenium) then leaving your test code in a separate repository is a fine way to go.

Chapter 6

Writing Your First Test

Fundamentally, Selenium works with two pieces of information, the element on a page you want to use and what you want to do with it. This one-two punch will be repeated over and over until you achieve the outcome you want in your application. At which point you will perform an assertion to confirm that the result is what you intended.

Let's take logging in to a website as an example. With Selenium you would:

1. Visit the page with the login form
2. Find the login form's username field and input text
3. Find the login form's password field and input text
4. Find the login form submit button and click it

Selenium is able to find and interact with elements on a page by way of various "locator strategies". The list includes `Class`, `CSS`, `ID`, `Link Text`, `Name`, `Partial Link Text`, `Tag Name`, and `XPath`.

While each serves a purpose, you only need to know a few to start writing effective tests.

How To Find Locators

The simplest way to find locators is to inspect the elements on a page. The best way to do this is from within your web browser. Fortunately, modern web browsers come pre-loaded with development tools that make this simple to accomplish.

When viewing the page, right-click on the element you want to interact with and click Inspect Element. This will bring up a small window with all of the HTML for the page but zoomed into your highlighted selection. From here you can see if there are unique or descriptive attributes you can work with.

How To Find Quality Elements

Your focus with picking an effective element should be on finding something that is unique, descriptive, and unlikely to change. Ripe candidates for this are `id` and `class` attributes. Whereas copy (e.g., text, or the text of a link) is less ideal since it is more apt to change. This may not hold true for when you make assertions, but it's a good goal to strive for.

If the elements you are attempting to work with don't have unique `id` or `class` attributes directly on them, look at the element that houses them (a.k.a. the parent element). Oftentimes the parent element has a unique locator that you can use to start with and drill down into the

element you want to use.

And if you can't find any unique elements, have a conversation with your development team letting them know what you are trying to accomplish. It's not hard for them to add helpful, semantic markup to make test automation easier. Especially when they know the use case you are trying to automate. The alternative can be a lengthy, painful process which will probably yield working test code, but it will be brittle and hard to maintain.

An Example

Part 1: Find The Elements And Write The Test

Here's the markup for a standard login form (pulled from [the login example on the internet](#)).

```
<form name="login" id="login" action="/authenticate" method="post">
  <div class="row">
    <div class="large-6 small-12 columns">
      <label for="username">Username</label>
      <input type="text" name="username" id="username">
    </div>
  </div>
  <div class="row">
    <div class="large-6 small-12 columns">
      <label for="password">Password</label>
      <input type="password" name="password" id="password">
    </div>
  </div>
  <button class="radius" type="submit"><i class="icon-2x icon-signin"> Login
</i></button>
</form>
```

Notice the unique elements on the form. The username input field has a unique `id` (e.g., `<input type="text" name="username" id="username">`), as does the password input field (e.g., `<input type="password" name="password" id="password">`). The submit button doesn't have an `id`, but it's the only button on the page.

Let's put these elements to use in our first test (or 'spec' as it's called in RSpec). First, let's make sure we have the correct versions of `rspec` and `selenium-webdriver` installed.

```
# filename: Gemfile
source 'https://rubygems.org'

gem 'rspec', '~> 3.4.0'
gem 'selenium-webdriver', '2.48.1'
```

After updating our Gemfile we can install the specified libraries with `bundle install` in the command-prompt. Now let's create our spec file.

```
# filename: login_spec.rb

require 'selenium-webdriver'

describe 'Login' do

  before(:each) do
    @driver = Selenium::WebDriver.for :firefox
  end

  after(:each) do
    @driver.quit
  end

  it 'succeeded' do
    @driver.get 'http://the-internet.herokuapp.com/login'
    @driver.find_element(id: 'username').send_keys('tomsmith')
    @driver.find_element(id: 'password').send_keys('SuperSecretPassword!')
    @driver.find_element(css: 'button').submit
  end

end
```

There are some new keywords in this code example that may not look familiar. Things like `describe`, `before(:each)`, `after(:each)`, and `it`. They are part of RSpec's syntax. Let's step through the file piece by piece. After that, it should all make sense.

At the top of the file we are pulling in `selenium-webdriver` which is the official open-source library for Selenium. It gives us the bindings necessary to drive the browser.

Next is the `describe`, which is the title of this test suite (or test class). A simple and helpful name is chosen to note the suite's intent (e.g., 'Login'). An opening `do` and a closing `end` (placed at the end of the file) are required to start and finish the suite.

In order to use Selenium we need to create an instance of it. This is done with the command `Selenium::WebDriver.for :firefox`, which is made available to us through the `selenium-webdriver` library. After our test we don't want the browser we opened to stick around, so we close it. This is what we're doing in `before(:each)` and `after(:each)`. Notice that we are storing our instance of Selenium in an instance variable (`@driver`) so we can access it throughout our entire test suite.

Last up is the `it` block, which is the test. Similar to `describe` we can give it a simple and helpful name (e.g., 'succeeded').

In the test we are using the locators we identified from the markup by finding them with `@driver.find_element`, passing in the locator type as a symbol (e.g., `id:`) and the locator value as a string (e.g., `username`, `password`, etc.). Once found, we take an action against each of them, which in this case is inputting text (with `.send_keys`) and click the submit button (with `.click`).

If we run this (e.g., `rspec login_spec.rb` from the command-line), it will run and pass. But there's one thing missing -- an assertion. Without it, there's no way to tell if the end state of our page is correct. We need to see what the page looks like after logging in, look through the markup, and find an element to use in our assertion.

Part 2: Figure Out What To Assert

Here is the markup that renders on the page after submitting the login form:

```
<div class="row">
  <div id="flash-messages" class="large-12 columns">
    <div data-alert="" id="flash" class="flash success">
      You logged into a secure area!
      <a href="#" class="close">x</a>
    </div>
  </div>
</div>

<div id="content" class="large-12 columns">
  <div class="example">
    <h2><i class="icon-lock"></i> Secure Area</h2>
    <h4 class="subheader">Welcome to the Secure Area. When you are done click logout
below.</h4>
    <a class="button secondary radius" href="/logout"><i class="icon-2x icon-signout">
Logout</i></a>
  </div>
</div>
```

After logging in, there looks to be a couple of things we can use for our assertion. There's the flash message class (most appealing), the logout button (appealing), or the copy from the h2 or the flash message (least appealing). Since the flash message class name is descriptive, denotes success, and is less likely to change than the copy, let's go with that.

But whenever we try to access a class element with two words in it (e.g., `class="flash success"`) we will need to combine them together into a single locator string that is written in either a CSS selector or XPath. Either approach works well, but the examples throughout this book will focus on how to use CSS selectors (when not using ID or Class locators).

A Quick Primer on CSS Selectors

In web design, CSS (Cascading Style Sheets) are used to apply styles to the markup (HTML) on a page. CSS is able to do this by declaring which bits of the markup it wants to interact with through the use of selectors. Selenium operates in a similar manner but instead of changing the style of elements, it interacts with them by clicking, getting values, typing, etc.

CSS selectors are a pretty straightforward and handy way to write locators, especially for hard to reach elements.

For right now, here's what you need to know. In CSS, class names start with a dot (`.`). For classes with multiple words, put a dot in front of each word, and remove the spaces (e.g., `.flash.success` for `class='flash success'`).

For a good resource on CSS Selectors, I encourage you to check out [Sauce Labs' write up on them](#).

Part 3: Write The Assertion And Verify It

```
# filename: login_spec.rb

require 'selenium-webdriver'

describe 'Login' do

  before(:each) do
    @driver = Selenium::WebDriver.for :firefox
  end

  after(:each) do
    @driver.quit
  end

  it 'succeeded' do
    @driver.get 'http://the-internet.herokuapp.com/login'
    @driver.find_element(id: 'username').send_keys('username')
    @driver.find_element(id: 'password').send_keys('password')
    @driver.find_element(css: 'button').click
    expect(@driver.find_element(css: '.flash.success').displayed?).to eq true
  end

end
```

Now our test has an assertion! But there's a lot going on in that one line, let's step through it to make sure we understand what it's doing.

First, we are finding the element we want to make an assertion against (using the locator type `css` and the locator string `.flash.success`). After that, we are asking Selenium if this element is displayed on the page (with `.displayed?`). This returns a boolean response that we can make an

assertion against. To make the assertion we leverage an RSpec assertion method (`expect`) and an RSpec matcher (`eq`) followed by the expected result `true`.

After logging in if the success flash message is displayed then Selenium will return `true` and the test will pass.

Just To Make Sure

Now when we run this test (`rspec login_spec.rb` from the command-line) it will pass just like before, but now there is an assertion which should catch a failure if something is amiss.

Just to make certain that this test is doing what we think it should, let's change the assertion to force a failure and run it again. A simple fudging of the locator will suffice.

```
expect(@driver.find_element(css: '.flash.successasdf').displayed?).to eq true
```

If it fails, then we can feel confident that it's doing what we expect, and can change the assertion back to normal before committing our code. This trick will save you more trouble than you know. Practice it often.

Chapter 7

Verifying Your Locators

If you're fortunate enough to be working with unique IDs and Classes, then you're usually all set. But when you have to handle more complex actions like traversing a page, or you need to run down odd test behavior it can be a real challenge to verify that you have the right locators to accomplish what you want.

Instead of the painful and tedious process of trying out various locators in your tests until you get what you're looking for, give Firefinder a try.

[Firefinder](#) is an add-on to the popular web-development Firefox tool [Firebug](#).

You first need to install Firebug, then Firefinder. Once you have that, verifying locators is a trivial task. And it works for both CSS and XPath locators.

NOTE: An alternative to FireFinder is [FirePath](#). If you have problems with FireFinder, or if you just want to see which one works better for you, then give it a try as well.

An Example

Let's try to identify the locators necessary to traverse a few levels into a large set of nested divs.

```
# a snippet from http://the-internet.herokuapp.com/large
```

```
<div id='siblings'>
  <div id='sibling-1.1'>1.1
  <div id='sibling-1.2'>1.2</div>
  <div id='sibling-1.3'>1.3</div>
  <div id='sibling-2.1'>2.1
  <div id='sibling-2.2'>2.2</div>
  <div id='sibling-2.2'>2.3</div>
  <div id='sibling-3.1'>3.1
  <div id='sibling-3.2'>3.2</div>
  <div id='sibling-3.2'>3.3</div>
  <div id='sibling-3.1'>4.1
  <div id='sibling-3.2'>4.2</div>
  <div id='sibling-3.2'>4.3</div>
  ...
```

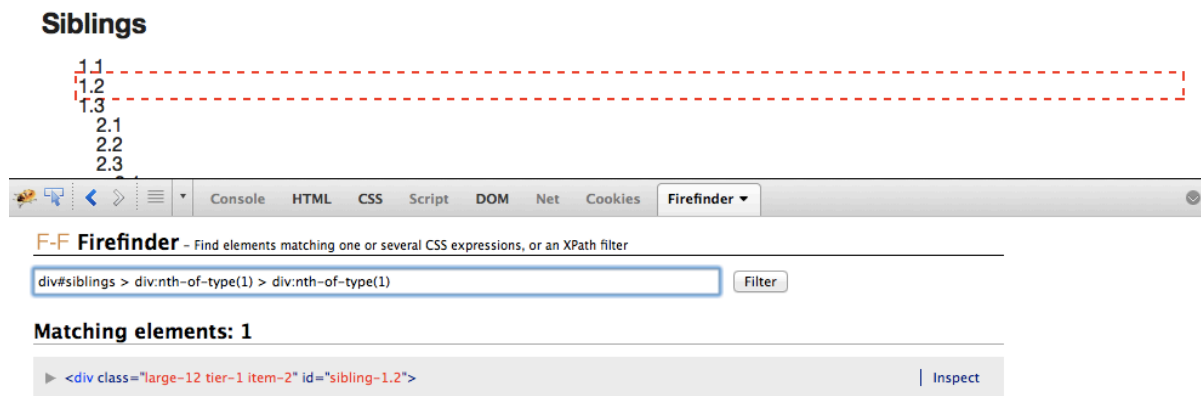
If we were to perform a `find_element` action using the following locator, it will work.

```
@driver.find_element(css: "div#siblings > div:nth-of-type(1) > div:nth-of-type(1)")
```

But if we try to go one level deeper, it won't work.

```
@driver.find_element(css: "div#siblings > div:nth-of-type(1) > div:nth-of-type(1) > div:nth-of-type(1)")
```

Fortunately with Firefinder we can actually see what our locators are doing. Here's what it shows us for the locators that "worked".



It looks like our locators are scoping to the wrong part of the first level (1.2). We need to reference the third part of each level (e.g., 1.3, 2.3, 3.3) in order to traverse deeper since the nested divs live under the third part of each level.

So if we try this locator instead, it should work.

```
@driver.find_element(css: "div#siblings > div:nth-of-type(1) > div:nth-of-type(3) > div:nth-of-type(3)")
```

And we can confirm this before changing any test code by looking in Firefinder.

Siblings

1.1

1.2

1.3

2.1

2.2

2.3

3.1

3.2

3.3

4.1

4.2

4.3

5.1

5.2

5.3

6.1

6.2

6.3

7.1

7.2

Firefinder

Find elements matching one or several CSS expressions, or an XPath filter

div#siblings > div:nth-of-type(1) > div:nth-of-type(3) > div:nth-of-type(3)

Filter

Matching elements: 1

<div class="parent large-12 columns tier-3 item-1" id="sibling-3.1">

Inspect

This should help save you time and frustration when running down tricky locators in your tests. It definitely has for me.

Chapter 8

Writing Re-usable Test Code

One of the biggest challenges with Selenium tests is that they can be brittle and challenging to maintain over time. This is largely due to the fact that things in the app you're testing change, breaking your tests.

But the reality of a software project is that change is a constant. So we need to account for this reality somehow in our test code in order to be successful.

Enter Page Objects.

Rather than write your test code directly against your app, you can model the behavior of your application into simple objects and write your tests against them instead. That way when your app changes and your tests break, you only have to update your test code in one place to fix it.

With this approach we not only get the benefit of controlled chaos, we also get the benefit of reusable functionality across our tests.

An Example

Part 1: Create A Page Object And Update Test

Let's take our login example from the previous chapter and pull it out into a page object and update our test accordingly.

```

# filename: login.rb

class Login

  USERNAME_INPUT = { id: 'username' }
  PASSWORD_INPUT = { id: 'password' }
  SUBMIT_BUTTON = { css: 'button' }
  SUCCESS_MESSAGE = { css: '.flash.success' }

  def initialize(driver)
    @driver = driver
    @driver.get 'http://the-internet.herokuapp.com/login'
  end

  def with(username, password)
    @driver.find_element(USERNAME_INPUT).send_keys(username)
    @driver.find_element(PASSWORD_INPUT).send_keys(password)
    @driver.find_element(SUBMIT_BUTTON).click
  end

  def success_message_present?
    @driver.find_element(SUCCESS_MESSAGE).displayed?
  end

end

```

We start by creating our own class, naming it accordingly (`Login`), and storing our locators in hashes within constants along the top.

We then use an `initialize` method. This is a built in method for classes in Ruby that enable you to automatically execute code upon instantiating the class (a.k.a. a constructor). With it we are taking an argument to receive the Selenium driver object and storing it in an instance variable. This will enable the class to drive the browser.

In our `with` method we are capturing the core behavior of the login page by accepting the username and password input values and putting them to use with the typing and clicking actions for the login form.

Since our behavior now lives in a page object, we want a clean way to make an assertion in our test. This is where `success_message_present?` comes in. Notice that it ends with a question mark. Methods that end like this are known as mutator methods. When they end with a question mark, they imply that they will return a boolean. So in it we want to ask a question of the page with Selenium so that we can return a boolean (e.g.,

`@driver.find_element(SUCCESS_MESSAGE).displayed?`). In our test we will put this boolean to work by making an assertion against it.

Now let's update our test to use this page object.

```
# filename: login_spec.rb

require 'selenium-webdriver'
require_relative 'login'

describe 'Login' do

  before(:each) do
    @driver = Selenium::WebDriver.for :firefox
    @login = Login.new(@driver)
  end

  after(:each) do
    @driver.quit
  end

  it 'succeeded' do
    @login.with('tomsmith', 'SuperSecretPassword!')
    expect(@login.success_message_present?).to eq true
  end

end
```

At the top of the file we include the page object by using `require_relative`. This enables us to reference another file based on the current file's path.

Next we instantiate our login page object in `before(:each)`, passing in `@driver` as an argument, and storing it all in an instance variable (`@login`). We then wire up our test to use the actions available in `@login`.

We finish things up by making an assertion against our helper method (`expect(@login.success_message_present?).to eq true`).

Part 2: Write Another Test

This may feel like more work than what we had when we first started. But we're in a much sturdier position and able to write follow-on tests more easily. Let's add another test to demonstrate a failed login.

If we provide incorrect credentials, the following markup gets rendered on the page.

```
<div id="flash-messages" class="large-12 columns">
  <div data-alert="" id="flash" class="flash error">
    Your username is invalid!
    <a href="#" class="close">x</a>
  </div>
</div>
```

This is similar to the markup from our successful flash message, so the mechanics will be similar in our page object. First we'll add a new locator for the failure message in our list of constants at the top of our class (just below our success message locator).

```
# filename: login.rb

class Login

  USERNAME_INPUT = { id: 'username' }
  PASSWORD_INPUT = { id: 'password' }
  SUBMIT_BUTTON  = { css: 'button' }
  SUCCESS_MESSAGE = { css: '.flash.success' }
  FAILURE_MESSAGE = { css: '.flash.error' }
  #...
```

Further down the file (just after the existing mutator method) we'll add a new method to check to see if the failure message is displayed.

```
# filename: login.rb
# ...
def success_message_present?
  @driver.find_element(SUCCESS_MESSAGE).displayed?
end

def failure_message_present?
  @driver.find_element(FAILURE_MESSAGE).displayed?
end

end
```

Lastly, we add a new test in our spec file just below our existing one, specifying invalid credentials to force a failure.

```
# filename: login_spec.rb
# ...
it 'succeeded' do
  @login.with('tomsmith', 'SuperSecretPassword!')
  expect(@login.success_message_present?).to eq true
end

it 'failed' do
  @login.with('asdf', 'asdf')
  expect(@login.failure_message_present?).to eq true
end

end
```

Now if we run our spec file (`rspec login_spec.rb`) we will see two browser windows open (one after the other) testing both the successful and failure login conditions.

Why Asserting False Won't Work (yet)

You may be wondering why we didn't check to see if the success message wasn't present.

```
it 'failed' do
  @login.with('asdf', 'asdf')
  expect(@login.success_message_present?).to eq false
end
```

There are two problems with this approach. First, the absence of success message does not necessarily indicate a failed login. The assertion we ended up with is more concise. Second, our test will fail. This is because it errors out when looking for an element that's not present.

```
.F
```

```
Failures:
```

```
1) Login failed
```

```
Failure/Error: expect(@login.success_message_present?).to eql false
```

```
Selenium::WebDriver::Error::NoSuchElementException:
```

```
Unable to locate element: {"method":"css selector","selector":".flash.success"}
```

```
# [remote server]
```

```
file:///var/folders/sk/gfclf5qjlvv_z4frjbvhr62w0000gn/T/webdriver-profile20...
```

```
# [remote server]
```

```
file:///var/folders/sk/gfclf5qjlvv_z4frjbvhr62w0000gn/T/webdriver-profile20...
```

```
# [remote server]
```

```
file:///var/folders/sk/gfclf5qjlvv_z4frjbvhr62w0000gn/T/webdriver-profile20...
```

```
# [remote server]
```

```
file:///var/folders/sk/gfclf5qjlvv_z4frjbvhr62w0000gn/T/webdriver-profile20...
```

```
# [remote server]
```

```
file:///var/folders/sk/gfclf5qjlvv_z4frjbvhr62w0000gn/T/webdriver-profile20...
```

```
# ./login.rb:21:in `success_message_present?'
```

```
# ./login_spec.rb:23:in `block (2 levels) in <top (required)>'
```

```
Finished in 8.85 seconds (files took 0.18961 seconds to load)
```

```
2 examples, 1 failure
```

```
Failed examples:
```

```
rspec ./login_spec.rb:20 # Login failed
```

We'll address this limitation in the next chapter.

Part 3: Confirm We're In The Right Place

Before we can call our page object complete, there's one more addition we'll want to make. We'll want to add a check in the constructor to make sure the page is in the correct state. Otherwise the test should not proceed and fail.

As a rule, you want to keep assertions in your tests and out of your page objects. So we'll use an exception instead.

```
# filename: login.rb

class Login

  LOGIN_FORM      = { id: 'login' }
  USERNAME_INPUT  = { id: 'username' }
  PASSWORD_INPUT  = { id: 'password' }
  SUBMIT_BUTTON   = { css: 'button' }
  SUCCESS_MESSAGE = { css: '.flash.success' }
  FAILURE_MESSAGE = { css: '.flash.error' }

  def initialize(driver)
    @driver = driver
    @driver.get 'http://the-internet.herokuapp.com/login'
    raise 'Login page not ready' unless
      @driver.find_element(LOGIN_FORM).displayed?
  end
  # ...
end
```

We add a new locator to the mix (the `id` of the login form) and add a new line to the end of our `initialize` method. In it we are checking to see if the login form is displayed. If it's not, we raise a custom exception.

If we run our tests again, they should pass just like before. But now we can rest assured that the test will only proceed if the login form is present.

Chapter 9

Writing Really Re-usable Test Code

In the previous chapter we stepped through creating a simple page object to capture the locators and behavior of the page we were interacting with. While this was a good start, it leaves a lot of room for improvement.

As our test suite grows and we add more page objects we will start to see common behavior that we will want to use over and over again -- across numerous page objects. If we leave this unchecked we will end up with duplicative code which will slowly make our test suite harder to maintain. Also, right now we're using Selenium actions directly in our page object. While on the face of it this may seem fine, it has some long term impacts:

- the inability to drive your tests with a different driver (e.g., mobile)
- test maintenance issues when Selenium changes (e.g., major upgrades ala RC to WebDriver)
- slower test writing due to the lack of a simple Domain Specific Language (DSL)

With a Base Page Object (a.k.a. a facade layer) we can easily side step these concerns by abstracting all of our common actions into a central location and leverage them in our page objects.

Let's step through an example with our login page object.

An Example

Part 1: Create The Base Page Object

First we will need to create the base page object.

```
# filename: base_page.rb
require 'selenium-webdriver'

class BasePage

  def initialize(driver)
    @driver = driver
  end

  def visit(url)
    @driver.get url
  end

  def find(locator)
    @driver.find_element locator
  end

  def type(text, locator)
    find(locator).send_keys text
  end

  def click(locator)
    find(locator).click
  end

  def is_displayed?(locator)
    find(locator).displayed?
  end

end
```

First we pull in the Selenium Ruby bindings (`require selenium-webdriver`). By placing it here in this file we will no longer need to require it anywhere else in our test suite (assuming all of our tests leverage page objects). And it makes sense to have it live here since this is the only file where we will be referencing Selenium commands directly.

Next we declare our class and name it appropriately with camel-casing (`class BasePage`).

In our `initialize` method we are accepting a driver object as an argument and setting it as an instance variable to make it available throughout this class. We then add common methods that we will reference in our page objects (e.g., `visit`, `find`, `type`, `submit`, and `is_displayed?`).

Now let's update our login page object to leverage this.

```

# filename: login.rb
require_relative 'base_page'

class Login < BasePage

  LOGIN_FORM      = { id: 'login' }
  USERNAME_INPUT  = { id: 'username' }
  PASSWORD_INPUT  = { id: 'password' }
  SUBMIT_BUTTON   = { css: 'button' }
  SUCCESS_MESSAGE = { css: '.flash.success' }
  FAILURE_MESSAGE = { css: '.flash.error' }

  def initialize(driver)
    super
    visit 'http://the-internet.herokuapp.com/login'
    raise 'Login page not ready' unless
      is_displayed?(LOGIN_FORM)
  end

  def with(username, password)
    type username, USERNAME_INPUT
    type password, PASSWORD_INPUT
    click SUBMIT_BUTTON
  end

  def success_message_present?
    is_displayed? SUCCESS_MESSAGE
  end

  def failure_message_present?
    is_displayed? FAILURE_MESSAGE
  end

end

```

In order to use the base page object we need to do two things. First we need to include the file, which is handled using `require_relative`. And second, we need to connect the login page object with the base page object with inheritance (which is handled in the class declaration with the `<` operator). This is effectively saying `Login` is a child of `BasePage`. Or, `BasePage` is the parent of `Login`. This enables us to freely reference the methods we created in the base page object.

As a result of this refactoring, our `initialize` method changes a bit. Rather than setting the driver object as an instance variable, we are calling `super`. This is a built-in keyword in Ruby for class inheritance. It triggers a method of the same name to be run from the parent class. So in this case, `initialize` from `BasePage` gets run, which is responsible for handling the driver object retrieval and instance variable creation.

Since the `BasePage` is now the only place where we reference Selenium actions directly, it stands that it should be the only place where we need to reference a `@driver` object. So we replace all of the driver incantations in this class with calls to the methods we created in the base page object. Our tests remain unchanged, but our page objects are now more flexible and readable.

Running our tests (`rspec login_spec.rb` from the command-line) will still yield a passing result.

Part 2: Add Some Error Handling

Remember in the previous chapter when we ran into an error with Selenium when we looked for an element that wasn't on the page? Let's address that now.

To recap -- here's the error message we saw:

```
.F

Failures:

  1) Login failed
     Failure/Error: expect(@login.success_message_present?).to eq false

Selenium::WebDriver::Error::NoSuchElementException:
  Unable to locate element: {"method":"css selector","selector":".flash.success"}
    # [remote server]
file:///var/folders/sk/gfclf5qjlvv_z4frjbvhr62w0000gn/T/webdriver-profile20...
    # [remote server]
file:///var/folders/sk/gfclf5qjlvv_z4frjbvhr62w0000gn/T/webdriver-profile20...
    # [remote server]
file:///var/folders/sk/gfclf5qjlvv_z4frjbvhr62w0000gn/T/webdriver-profile20...
    # [remote server]
file:///var/folders/sk/gfclf5qjlvv_z4frjbvhr62w0000gn/T/webdriver-profile20...
    # [remote server]
file:///var/folders/sk/gfclf5qjlvv_z4frjbvhr62w0000gn/T/webdriver-profile20...
    # ./login.rb:21:in `success_message_present?'
    # ./login_spec.rb:23:in `block (2 levels) in <top (required)>'

Finished in 8.85 seconds (files took 0.18961 seconds to load)
2 examples, 1 failure

Failed examples:

rspec ./login_spec.rb:20 # Login failed
```

The important thing to note is the specific error message Selenium offered up. The part that comes just before the line that says 'Unable to locate element'. We're interested in

`Selenium::WebDriver::Error::NoSuchElementException`. With it we can make our code catch it and return a `false` boolean instead of failing the test.

Let's update the `is_displayed?` method in our base page object to do just that.

```
# filename: base_page.rb
#...

def is_displayed?(locator)
  begin
    find(locator).displayed?
  rescue Selenium::WebDriver::Error::NoSuchElementException
    false
  end
end
```

In Ruby we can make our code more resilient by wrapping an action in a `begin / rescue` block. In this case we're wrapping our attempt to see if an element is displayed to the user. If the element is not on the page Selenium will raise a `NoSuchElementException` exception. When that happens, our code will now `rescue` for that specific error condition, and return `false` instead of failing the test.

Now we can write a test to see if an element is not on the page without fear of our test blowing up.

Just to verify, let's revisit our failed login test and update it to assert that the success message isn't present.

```
# filename: login_spec.rb
#...

it 'failed' do
  @login.with('asdf', 'asdf')
  #expect(@login.failure_message_present?).to eq true
  expect(@login.success_message_present?).to eq false
end

end
```

Now when we run this (`rspec login_spec.rb` from the command-line) it will run without error, and pass. Feel free to change the assertion back when you're done.

Chapter 10

Writing Resilient Test Code

Ideally, you should be able to write your tests once and run them across all supported browsers. While this is a rosy proposition, there is some work to make this a reliable success. And sometimes there may be a hack or two involved. But the lengths you must go really depends on the browsers you care about and the functionality you're dealing with in your application.

By using high quality locators we're already in good shape, but there are still some issues to deal with. Most notably... timing. This is especially true when working with dynamic, JavaScript heavy pages (which is more the rule than the exception in a majority of web applications you'll deal with).

But there is a simple approach that makes up the bedrock of reliable and resilient Selenium tests -- and that's how you wait and interact with elements. Gone are the days of waiting for the page to finish loading, or hard-coding sleeps, or doing a blanket wait time (a.k.a. an implicit wait). Nay. Now are the wonder years of waiting for an expected outcome to occur for a set amount of time. If the outcome occurs before the amount of time specified, then the test will proceed. Otherwise, it will wait the full amount of time specified.

We accomplish this through the use of explicit waits.

An Example

Let's step through an example that demonstrates this against [a dynamic page on the-internet](#). The functionality is pretty simple -- there is a button. When you click it a loading bar appears for 5 seconds, then disappears, and gets replaced with the text 'Hello World!'.

Part 1: Create A New Page Object And Update The Base Page Object

Let's start by looking at the markup on the page.

```
<div class="example">
  <h3>Dynamically Loaded Page Elements</h3>
  <h4>Example 1: Element on page that is hidden</h4>

  <br>

  <div id="start">
    <button>Start</button>
  </div>

  <div id="finish" style="display:none">
    <h4>Hello World!</h4>
  </div>

</div>
```

At a glance it's simple enough to tell that there are unique `id` attributes that we can use to reference the start button and finish text.

Let's add a page object for Dynamic Loading.

```
# filename: dynamic_loading.rb
require_relative 'base_page'

class DynamicLoading < BasePage

  START_BUTTON = { css: '#start button' }
  FINISH_TEXT  = { id: 'finish' }

  def initialize(driver)
    super
    visit 'http://the-internet.herokuapp.com/dynamic_loading/1'
  end

  def start
    click START_BUTTON
  end

  def finish_text_present?
    wait_for(10) { is_displayed? FINISH_TEXT }
  end

end
```

At the top of the file we require our base page object and set up inheritance when declaring our

class so we get our common Selenium actions. After that we wire up our locators in constants, add a visit action to our `initialize` method, and drop in some methods to start the loading process and to see if the finish text is present.

In `finish_text_present?` we are using a method which has not been added to our base page object yet. So let's hop into our base page object and add it to the bottom of the class.

```
# filename: base_page.rb
# ...
def wait_for(seconds = 15)
  Selenium::WebDriver::Wait.new(timeout: seconds).until { yield }
end

end
```

`wait_for` is where we are defining our explicit wait. In it we are accepting an integer for the number of seconds we'd like to wait. If nothing is provided, 15 seconds will be used. We then use this value to tell the explicit wait function that Selenium offers how long to wait, and what to wait for. By using a `yield` we are able to easily pass in a code block to run (e.g., a Selenium command surrounded by brackets (`{ }`)).

More On Explicit Waits

It's important to set a reasonably sized default timeout for the explicit wait. But you want to be careful not to make it too high. Otherwise you run into a lot of the same timing issues you get from implicit waits. But set it too low and your tests will be brittle, forcing you to run down transient issues.

In our page object when we're using `wait_for(10) { is_displayed? FINISH_TEXT }` we are telling Selenium to see if the finish text is displayed on the page. It will keep trying until it either returns `true` or reaches ten seconds -- whichever comes first.

If the behavior on the page takes longer than we expect (e.g., due to slow load times, or a feature change), we can simply adjust this one wait time to fix the test rather than increase a blanket wait time (which impacts every test). And since it's dynamic, it won't always take the full amount of time to complete.

Part 2: Write A Test To Use The New Page Object

Now that we have our page object and requisite base page methods we can wire up a new test to use it.


```

# filename: dynamic_loading_spec.rb
require_relative 'dynamic_loading'

describe 'Dynamic Loading' do

  before(:each) do
    @driver = Selenium::WebDriver.for :firefox
    @dynamic_loading = DynamicLoading.new(@driver)
  end

  after(:each) do
    @driver.quit
  end

  it 'Example 1: Hidden Element' do
    @dynamic_loading.start
    expect(@dynamic_loading.finish_text_present?).to eql true
  end

end

```

When we run it (e.g., `rspec dynamic_loading_page.rb` from the command-line) it will visit the page, click the start button, wait for the text to appear, assert the text appeared, and close the browser.

Part 3: Update Page Object And Add A New Test

Let's step through one more dynamic page example to see if our explicit wait approach holds up.

[This example](#) is laid out similarly to the last one, the main difference is that it will render the final result after the progress bar completes. Here's the markup for it.

```

<div class="example">
  <h3>Dynamically Loaded Page Elements</h3>
  <h4>Example 2: Element rendered after the fact</h4>

  <br>

  <div id="start">
    <button>Start</button>
  </div>

  <br>
</div>

```

In order to find the selector for the finish text element we need to inspect the page after the loading bar sequence finishes. Here's what it looks like.

```
<div id="finish" style=""><h4>Hello World!</h4></div>
```

Before we add our test we need to modify our page object to accommodate visiting the different example URLs.

```
# filename: dynamic_loading.rb
require_relative 'base_page'

class DynamicLoading < BasePage

  START_BUTTON = { css: '#start button' }
  FINISH_TEXT = { id: 'finish' }

  def example(example_number)
    visit "http://the-internet.herokuapp.com/dynamic_loading/#{example_number}"
  end

  def start
    click START_BUTTON
  end

  def finish_text_present?
    wait_for(10) { is_displayed? FINISH_TEXT }
  end

end
```

We're able to remove the `initialize` method entirely since we no longer need the `visit` action as part of loading the page object. That functionality now lives in the `example` method. Thanks to inheritance our parent class is still able to do what it needs, even without the use of an `initialize` in the child class.

Now that we have that sorted, let's add a new test to reference the markup shown above (and update our existing test to use the new `.example` method).

```

# filename: dynamic_loading_spec.rb
require_relative 'dynamic_loading'

describe 'Dynamic Loading' do

  before(:each) do
    @driver = Selenium::WebDriver.for :firefox
    @dynamic_loading = DynamicLoading.new(@driver)
  end

  after(:each) do
    @driver.quit
  end

  it 'Example 1: Hidden Element' do
    @dynamic_loading.example 1
    @dynamic_loading.start
    expect(@dynamic_loading.finish_text_present?).to eq true
  end

  it 'Example 2: Rendered after the fact' do
    @dynamic_loading.example 2
    @dynamic_loading.start
    expect(@dynamic_loading.finish_text_present?).to eq true
  end

end

```

When we run these tests (`rspec dynamic_loading_spec.rb` from the command-line) we can see that the explicit wait approach works for when the element is on the page but hidden and when it's rendered after the fact.

Browser Timing

Using explicit waits gets you pretty far. But there are a few things you'll want to think about when it comes to writing your tests to work against various browsers.

It's simple enough to write your tests locally against Firefox and assume you're all set. But once you start to run things against other browsers, you may be in for a rude awakening. The first thing you're likely to run into is the speed of execution. A lot of your tests will start to fail when you point them at either Chrome or Internet Explorer, and likely for different reasons.

Chrome execution is very fast, so you will see some odd timeout failures. This is an indicator that you need to add explicit waits to parts of your page objects that don't already have them. And the inverse is true when running things against older version of Internet Explorer (e.g., IE 8). This is an

indicator that your explicit wait times are not long enough since the browser is taking longer to respond -- so your tests timeout.

The best approach to solve this is an iterative one. Pick a browser you care about and run your tests. Take each failed test, adjust your test code as needed, and run them again. Keep going until they all pass. Repeat for each of the browsers you care about.

Once you're on the other side of these issues, the amount of effort you need to put into it should diminish dramatically.

Closing Thoughts

By explicitly waiting to complete an action, our tests are in a much more resilient position because Selenium will keep trying for a reasonable amount of time rather than trying just once. And each action can be tuned to meet the needs of each circumstance. Couple that with the dynamic nature of explicit waits, and you have something that will work in a multitude of circumstances -- helping you endure even the toughest of browsers.

This is one of the most important concepts in testing with Selenium. Use explicit waits often.

Chapter 11

Prepping For Use

Now that we have tests, page objects, and a base page object, let's package things into a more useful structure.

Folder Organization

It's about time we create some folders for our tests (specs) and page objects. To err on the side of simplicity, let's call the folders 'spec' (for our tests) and 'pages' (for our page objects). We are using `spec` (singular) since it's a default folder that RSpec will know to look for.

Here's everything we should have after creating folders and moving files around:

```
.
|-- Gemfile
|-- pages
|   |-- base_page.rb
|   |-- dynamic_loading.rb
|   `-- login.rb
`-- spec
    |-- dynamic_loading_spec.rb
    |-- login_spec.rb
```

Updating Require Statements

As a result of doing this, we will need to update the require statements in our tests.

```
# filename: spec/login_spec.rb

require_relative '../pages/login'

describe 'Login' do
  ...
```

```
# filename: spec/dynamic_loading_spec.rb

require_relative '../pages/dynamic_loading'

describe 'Dynamic Loading' do
  ...
```

Notice the use of double-dots (`..`) in the page object require statement. This is how we tell Ruby to traverse up a directory. So `../pages/dynamic_loading` is traversing up from the spec directory, into the page objects folder, and then accessing `dynamic_loading`.

Base URL

Up until now we've been hard-coding the URL we want to use for our application. But it's likely that the application could have numerous URL end-points. In the case of [the-internet](http://the-internet.herokuapp.com) there is the production URL (e.g., `http://the-internet.herokuapp.com`) and then there's running the app locally (e.g., `http://localhost:4567`).

In order to service both of these URLs with our tests we'll want to make it so we can specify a base URL at runtime and update our page objects to use it. And we want to set a sensible default for the base URL so if we specify nothing, a valid URL will be provided to our tests. So let's create a central file to store configuration values like this and make sure it can receive values at runtime. Let's create a `config.rb` file in the parent directory and place it there.

```
# filename: config.rb

ENV['base_url'] ||= 'http://the-internet.herokuapp.com'
```

By using a conditional when setting the environment variable (`||=`) gets us our sensible default. We are making it so we can override this value when launching our test suite. It essentially means if the environment variable already exists and contains a value, use it. Otherwise, set it to 'http://the-internet.herokuapp.com'. This will come in handy later, and is something we'll use frequently.

With base URL in a central place we can now go and update our page objects to use it.

```
# filename: pages/base_page.rb
require 'selenium-webdriver'

class BasePage

  def initialize(driver)
    @driver = driver
  end

  def visit(url_path)
    @driver.get ENV['base_url'] + url_path
  end

  # ...
```

In our base page object we put the base URL environment variable to use in `visit`. By doing this

we update the argument for the method from expecting an entire `url` to just the `url_path`. Now to update the use of `visit` in our page objects.

```
# filename: pages/dynamic_loading.rb
require_relative 'base_page'

class DynamicLoading < BasePage

  START_BUTTON = { css: '#start button' }
  FINISH_TEXT  = { id: 'finish' }

  def example(example_number)
    visit "/dynamic_loading/#{example_number}"
  end
end

# ...
```

```
# filename: pages/login.rb
require_relative 'base_page'

class Login < BasePage

  LOGIN_FORM      = { id: 'login' }
  USERNAME_INPUT  = { id: 'username' }
  PASSWORD_INPUT  = { id: 'password' }
  SUBMIT_BUTTON   = { css: 'button' }
  SUCCESS_MESSAGE = { css: '.flash.success' }
  FAILURE_MESSAGE = { css: '.flash.error' }

  def initialize(driver)
    super
    visit '/login'
    raise 'Login page not ready' unless
      is_displayed?(LOGIN_FORM)
  end
end

# ...
```

Now our page objects are free from hard-coded URLs, making our tests more flexible.

Global Setup And Teardown

Now we're ready to pull the test setup and teardown actions out of our tests and into a central place. In RSpec this is straight-forward through the use of a `spec_helper` file. This can live alongside the other test files in the `spec` directory.

```

# filename: spec/spec_helper.rb
require 'selenium-webdriver'

RSpec.configure do |config|

  config.before(:each) do
    @driver = Selenium::WebDriver.for :firefox
  end

  config.after(:each) do
    @driver.quit
  end

end

```

We require the Selenium library here since we're working directly with it. By having it here we can remove it from our base page object. We can also remove our Selenium commands from the `before(:each)` and `after(:each)` in our tests -- replacing them with a simple require statement at the top of the file (`require_relative 'spec_helper'`). We will also be able to remove the `after(:each)` method from our tests, leaving just the `before(:each)` for use with our page objects.

Here's what our tests look like with these changes:

```

# filename: spec/login_spec.rb
require_relative 'spec_helper'
require_relative '../pages/login'

describe 'Login' do

  before(:each) do
    @login = Login.new(@driver)
  end

  it 'succeeded' do
    @login.with('tomsmith', 'SuperSecretPassword!')
    expect(@login.success_message_present?).to eq true
  end

  it 'failed' do
    @login.with('asdf', 'asdf')
    expect(@login.failure_message_present?).to eq true
  end

end

```



```

# filename: spec/dynamic_loading_spec.rb
require_relative 'spec_helper'
require_relative '../pages/dynamic_loading'

describe 'Dynamic Loading' do

  before(:each) do
    @dynamic_loading = DynamicLoading.new(@driver)
  end

  it 'Example 1: Hidden Element' do
    @dynamic_loading.example 1
    @dynamic_loading.start
    expect(@dynamic_loading.finish_text_present?).to eq true
  end

  it 'Example 2: Rendered after the fact' do
    @dynamic_loading.example 2
    @dynamic_loading.start
    expect(@dynamic_loading.finish_text_present?).to eq true
  end

end

```

Running Everything

Now that things are cleaned up, let's run everything. To do that we'll want to make sure to include our new config file. We can do that by specifying it at run time with `rspec --require ./config.rb`, or, `rspec -r ./config.rb` (for short). Note the `./` before `config.rb`. This tells RSpec that the config file is in the current directory.

To specify a different base URL, prepend the command with it (e.g., `base_url=http://localhost:4567 rspec -r ./config.rb` if you had a local instance of [the-internet](#) running).

Give it a shot. All of the tests should run and pass.

Chapter 12

Running A Different Browser Locally

It's easy to get your tests running locally against Firefox since Selenium comes with it built-in. But when you want to run them against a different browser like Chrome you quickly run into configuration overhead that can seem overly complex and lacking in good documentation or examples.

A Brief Primer On Browser Drivers

With the introduction of WebDriver (circa Selenium 2) a lot of benefits were realized (e.g., more effective and faster browser execution, no more single host origin issues, etc). But with it came some architectural and configuration differences that may not be widely known. Namely -- browser drivers.

WebDriver works with each of the major browsers through a browser driver which is (ideally) maintained by the browser manufacturer. It is an executable file (consider it a thin layer or a shim) that acts as a bridge between Selenium and the browser.

Let's step through an example using [ChromeDriver](#).

An Example

Before starting we'll need to [grab the latest ChromeDriver binary executable from Google](#) and store the unzipped contents of it in our test code. The simplest thing to do is create a new folder for it (and other things like it).

Let's create a `vendor` directory and place it there.

```
.
|-- config.rb
|-- Gemfile
|-- pages
|   |-- base_page.rb
|   |-- dynamic_loading.rb
|   `-- login.rb
|-- spec
|   |-- dynamic_loading_spec.rb
|   |-- login_spec.rb
|   `-- spec_helper.rb
`-- vendor
    `-- chromedriver
```

In order for Selenium to use this binary, we have to make sure it knows where it is. There are two ways to do that. We can add the `chromedriver` file to our system path, or pass in the path to the file when launching Selenium. Let's do the latter.

NOTE: There is a different chromedriver binary for each major operating system. If you're using Windows, be sure to use the one that ends in `.exe` and specify `chromedriver.exe` in your configuration. This example was built to run on OSX (which does not have a specified file extension for chromedriver).

But before we go too far, we want to make it so our test suite can run either Firefox or Chrome. To do that, we need to add a browser environment variable to our `'config.rb'` and some conditional logic to `spec_helper.rb`.

```
# filename: config.rb
```

```
ENV['base_url'] ||= 'http://the-internet.herokuapp.com'
ENV['browser']   ||= 'firefox'
```

```
# filename: spec/spec_helper.rb
```

```
require 'selenium-webdriver'
```

```
RSpec.configure do |config|
```

```
  config.before(:each) do
```

```
    case ENV['browser']
```

```
    when 'firefox'
```

```
      @driver = Selenium::WebDriver.for :firefox
```

```
    when 'chrome'
```

```
      Selenium::WebDriver::Chrome::Service.executable_path = File.join(Dir.pwd,
'vendor/chromedriver')
```

```
      @driver = Selenium::WebDriver.for :chrome
```

```
    end
```

```
  end
```

```
  config.after(:each) do
```

```
    @driver.quit
```

```
  end
```

```
end
```

In `config.before(:each)` we're doing a conditional check against the browser environment variable we set in `config.rb` to determine which browser driver to load. For Chrome we are first telling Selenium where the ChromeDriver binary executable is located on disk, and then loading an instance of Selenium for Chrome.

Assuming Chrome is already installed on the system (in a standard location) when we run our tests (specifying Chrome as the browser) -- `browser='chrome' rspec -r ./config.rb`, then we should see them execute in Chrome. If that's not the case for you, then check out [the ChromeDriver documentation site](#) to see if there are additional requirements for your specific system setup.

It's worth noting that this will only be reasonably performant since it is launching and terminating the ChromeDriver binary executable before and after every test. There are alternative approaches we can take, but this is good enough to see where our tests fall down in Chrome (and it will not be the primary way we will run our tests a majority of the time).

Additional Browsers

A similar approach can be applied to other browser drivers, with the only real limitation being the operating system you're running. But remember, no two browser drivers are alike. Be sure to check out the documentation for the browser you care about to find out it's specific requirements:

- [ChromeDriver](#)
- [FirefoxDriver](#)
- [InternetExplorer Driver](#)
- [SafariDriver](#)

Chapter 13

Running Browsers In The Cloud

If you've ever needed to test features in an older browser like Internet Explorer 8 then odds are you ran a virtual machine (VM) on your computer with a "legit" version of Windows XP.

Handy, but what happens when you need to check things on multiple versions of IE? Now you're looking at multiple VMs. And what about when you need to scale and cover other browser and Operating System (OS) combinations? Now you're looking at provisioning, running, and maintaining your own farm of machines and standing up something like Selenium Grid to coordinate tests across them.

And all you wanted to do was run your tests on the browsers you cared about...

Rather than take on the overhead of a test infrastructure you can easily outsource things to a third-party cloud provider. There are a handful of players in this space, but there's one that stands out -- Sauce Labs.

A Selenium Remote, Selenium Grid, And Sauce Labs Primer

At the heart of Selenium at scale is the use of Selenium Grid and Selenium Remote. Selenium Grid lets you distribute test execution across several machines and you connect to it with Selenium Remote. With Selenium Remote you specify the browser type, browser version, and operating system through the use of Selenium Remote's `Capabilities`.

This is fundamentally how Sauce Labs works. Behind the curtain they are ultimately running Selenium Grid, and they receive and execute your tests through Selenium Remote.

Let's dig in with an example.

An Example

Part 1: Initial Setup

NOTE: You'll need an account to use Sauce Labs. Their free trial offers more than enough to get you started ([link](#)).

With Sauce Labs we need to provide our credentials, specifics about what we want in our test environment, and configure Selenium a little bit differently than we have been. Let's start by creating a new config file for cloud execution.

```
# filename: config_cloud.rb

ENV['base_url']      ||= 'http://the-internet.herokuapp.com'
ENV['host']           = 'saucelabs'
# ENV['SAUCE_USERNAME'] = 'your-sauce-username'
# ENV['SAUCE_ACCESS_KEY'] = 'your-sauce-access-key'
ENV['operating_system'] ||= 'Windows XP'
ENV['browser']         ||= 'internet_explorer'
ENV['browser_version'] ||= '8'
```

Notice the use of a host environment variable (e.g., `ENV['host'] = 'saucelabs'`). This is what we'll use in our `spec_helper` file to determine whether to run things locally or in the cloud. We can specify our Sauce Labs credentials here by hard-coding them, or we configure them on our system. And we'll use the `operating_system`, `browser`, and `browser_version` environment variables to populate the `Capabilities` object.

```

# filename: spec/spec_helper.rb
require 'selenium-webdriver'

RSpec.configure do |config|

  config.before(:each) do |example|
    case ENV['host']
    when 'saucelabs'
      caps = Selenium::WebDriver::Remote::Capabilities.send(ENV['browser'])
      caps.version = ENV['browser_version']
      caps.platform = ENV['operating_system']
      caps[:name] = example.metadata[:full_description]

      @driver = Selenium::WebDriver.for(
        :remote,
        url: "http://#{ENV['SAUCE_USERNAME']} :#{ENV['SAUCE_ACCESS_KEY']}
@ondemand.saucelabs.com:80/wd/hub",
        desired_capabilities: caps)
    else
      case ENV['browser']
      when 'firefox'
        @driver = Selenium::WebDriver.for :firefox
      when 'chrome'
        Selenium::WebDriver::Chrome::Service.executable_path = File.join(Dir.pwd,
'vendor/chromedriver')
        @driver = Selenium::WebDriver.for :chrome
      end
    end
  end

  config.after(:each) do
    @driver.quit
  end
end
end

```

We've taken our browser conditional and made it nested underneath one for the host environment variable. If the host is set to `'saucelabs'`, then we configure the capabilities for Selenium Remote, passing in the requisite information that we will need for our Sauce Labs session. Otherwise, our tests will run locally.

There are a few things in this example that may be worth elaborating on.

First, we are using something called metaprogramming (a.k.a. code that writes code) when we are calling `Selenium::WebDriver::Remote::Capabilities`. We are using the `.send` method to pass in the environment variable. The value of which, in this case, is the same name as the method to

configure Selenium Remote to use Internet Explorer. So, we are in effect, specifying `Selenium::WebDriver::Remote::Capabilities.internet_explorer`. And if we were to specify 'chrome' for `ENV['browser']`, then it would give us `Selenium::WebDriver::Remote::Capabilities.chrome`.

Second, for `caps[:name]` we are using a piece of functionality built into RSpec which gives us the name of each test as it is being run. This will make it so each individual job in Sauce Labs will have the name of the test that was run.

Third, for the `url:` line in our `@driver` instantiation, we are injecting our environment variable through the use of string interpolation. This is why we are using double-quoted strings. If they were single-quotes then we wouldn't be able to do it.

Now if we run our test suite (`rspec -r ./config_cloud.rb`) and navigate to [our Sauce Labs Account page](#) then we should see each of the tests running in their own job, with proper names, against Internet Explorer 8.

Part 2: Test Status

The only thing missing now is the pass/fail status of the job. In our local terminal window everything should be coming up green. But in the Sauce Labs dashboard each of the test jobs will just say Finished. This will make our results less useful in the long run, so let's fix it.

Thanks to Sauce Labs' `sauce_whisk` gem it's simple to do.

```
# filename: Gemfile

source 'https://rubygems.org'

gem 'rspec', '~> 3.4.0'
gem 'selenium-webdriver', '2.48.1'
gem 'sauce_whisk', '~> 0.0.19'
```

After we install it we will need to require it somewhere, and `config_cloud.rb` is the most logical place since we will only need it when running our tests in Sauce. So let's add it to the top of the file.

```
# filename: config_cloud.rb

require 'sauce_whisk'

...
```

All that's left is to add an action to our `after(:each)` block in `spec_helper.rb` to use `sauce_whisk` to set the correct job status.

Before we issue `@driver.quit` we will want to grab the session ID from our `@driver` object (which is the job ID in Sauce Labs) and use it to set the status of the job based on the test result. Also, we'll want to make sure that it only executes when running tests against Sauce Labs, so we'll wrap it in a conditional check against the host environment variable.

```
# filename: spec/spec_helper.rb
#...

config.after(:each) do |example|
  if ENV['host'] == 'saucelabs'
    if example.exception.nil?
      SauceWhisk::Jobs.pass_job @driver.session_id
    else
      SauceWhisk::Jobs.fail_job @driver.session_id
    end
  end

  @driver.quit
end

end
```

Now when we run our tests (`rspec -r ./config_cloud.rb`) and navigate to [our Sauce Labs Account page](#), we should see our tests running like before. But now when they finish there should be a proper test status (e.g., 'Pass' or 'Fail').

Part 3: Job URL

If we leave things as they are we'll have a hard time correlating test failures with the Sauce Labs job that they ran in. Let's remedy that now.

```

# filename: spec/spec_helper.rb
# ...
config.after(:each) do |example|
  begin
    if ENV['host'] == 'saucelabs'
      if example.exception.nil?
        SauceWhisk::Jobs.pass_job @driver.session_id
      else
        SauceWhisk::Jobs.fail_job @driver.session_id
        raise "Watch a video of the test at https://saucelabs.com/tests/#{@driver.
session_id}"
      end
    end
  end
  ensure
    @driver.quit
  end
end
end

```

In `config.after(:each)` we `raise` a custom exception when there's a test failure and when we're running on Sauce Labs. The custom exception includes a URL to the Sauce Labs job which we create with `@driver.session_id`. In order for `@driver.quit` to run when an exception is raised, we wrap it with an `ensure` (inside of a `begin` block). This ensures that `@driver.quit` will execute regardless of what happens.

Now when we run our test suite and a test fails, it will read with a helpful URL to the Sauce Labs job.

Watch a video of the test at
<https://saucelabs.com/tests/8b075a49cb20477f9aa820de4d196ac5>

Part 4: Accessing Private Apps

There are various ways that companies make their pre-production application available for testing. Some use an obscure public URL and protect it with some form of authentication (e.g., Basic Auth, or cert based authentication). Others keep it behind their firewall. For those that stay behind a firewall, Sauce Labs has you covered.

They have a program called [Sauce Connect](#) that creates a secure tunnel between your machine and their cloud. With it, you can run tests in Sauce Labs and test applications that are only available on your private network.

To use Sauce Connect, you need to download and run it. There's a copy for each operating system -- get yours [here](#) and run it from the command-line. In the context of our existing test code let's

download Sauce Connect, unzip it's contents, and store it in a sub-folder in the `vendor` directory.

```
.
Gemfile
Gemfile.lock
config
  cloud.rb
  local.rb
pages
  base_page.rb
  dynamic_loading.rb
  login.rb
spec
  dynamic_loading_spec.rb
  login_spec.rb
  spec_helper.rb
vendor
  chromedriver
  sauce-connect
    bin
      sc
      sc.dSYM
        Contents
          Info.plist
          Resources
            DWARF
              sc
    include
      sauceconnect.h
    lib
      libsauceconnect.a
      libsauceconnect.la
  license.html
```

Now we just need to launch the application while specifying our Sauce account credentials.

```
> vendor/sauce-connect/bin/sc -u your-sauce-username -k your-sauce-access-key
Sauce Connect 4.3.11, build 1760 53519d5
Starting up; pid 38259
Command line arguments: vendor/sauce-connect/bin/sc -u the-internet -k ****
Using no proxy for connecting to Sauce Labs REST API.
Resolving saucelabs.com to 162.222.75.243 took 36 ms.
Started scproxy on port 64531.
Please wait for 'you may start your tests' to start your tests.
Starting secure remote tunnel VM...
Secure remote tunnel VM provisioned.
Tunnel ID: 7df798c046d042b2a90ef46db205293c
Secure remote tunnel VM is now: booting
Secure remote tunnel VM is now: running
Using no proxy for connecting to tunnel VM.
Resolving tunnel hostname to 162.222.75.23 took 43ms.
Starting Selenium listener...
Establishing secure TLS connection to tunnel...
Selenium listener started on port 4445.
Sauce Connect is up, you may start your tests.
```

Now that the tunnel is established, we could run our tests against a local instance of our application (e.g., [the-internet](#)). Assuming the application was set up and running on our local machine, we could run `base_url=http://localhost:4567 rspec -r ./config_cloud.rb` from the command-line and it would work.

To see the status of the tunnel, we can view it on [the tunnel page of the account dashboard](#). To shut the tunnel down, we can do it manually from this page. Or we can issue a `Ctrl+C` command to the terminal window where its running.

When the tunnel is closing, here's what you'll see.

```
Cleaning up.
Checking domain overlap for my domain sauce-connect.proxy, other tunnel domain sauce-
connect.proxy
Overlapping domain: sauce-connect.proxy, shutting down tunnel 7
df798c046d042b2a90ef46db205293c.
Goodbye.
```

Part 5: A Small Bit Of Clean-up

Now that we have two config files, let's move them into their own directory and pare down the names for simplicity. Let's create a `config` directory and rename the files to `local.rb` and `cloud.rb`.

This leaves our directory tree structure looking like this:

```
.
Gemfile
Gemfile.lock
config
  cloud.rb
  local.rb
pages
  base_page.rb
  dynamic_loading.rb
  login.rb
spec
  dynamic_loading_spec.rb
  login_spec.rb
  spec_helper.rb
vendor
  chromedriver
  sauce-connect
    bin
      sc
      sc.dSYM
        Contents
          Info.plist
          Resources
            DWARF
              SC
    include
      sauceconnect.h
    lib
      libsauceconnect.a
      libsauceconnect.la
    license.html
```

Chapter 14

Speeding Up Your Test Runs

It's great that we can easily run our tests in Sauce Labs. But it's a real bummer that all of our tests are executing in series. As our suite grows it will take longer and longer for it to finish running, which puts a real damper on our ability to get feedback quickly.

With parallelization we can remedy this. And in Ruby there is a library that makes this simple to accomplish.

Enter [parallel tests](#).

Setup

After installing the library we just have to change out the command we use to execute our test runs. Instead of `rspec` we will use one that `paralleltests` provides -- `parallelrspec`.

```
# filename: Gemfile
source 'https://rubygems.org'

gem 'rspec', '~> 3.4.0'
gem 'selenium-webdriver', '2.48.1'
gem 'sauce_whisk', '~> 0.0.19'
gem 'parallel_tests', '~> 2.0.0'
```

`parallel_rspec` is effectively a wrapper around RSpec. It's responsible for breaking our spec files into groups and launching each of them in separate system processes along with any arguments we pass in to configure RSpec. So we can still provide our `-r` for a config file, we just have to specify `--test-options` before-hand.

To run our tests in parallel with Sauce Labs (without the secure tunnel), then we would use `parallel_rspec --test-options '-r ./config/cloud.rb' spec`. Note the additional `spec` at the end of the command. This tells `parallel_rspec` where the test files live (which is in the `spec` directory).

These extra arguments can seem a little verbose, but it's a small price to pay for gaining parallel test execution.

Randomizing

A great way to make sure your tests don't have any inter-dependencies (and to ferret out possible anomalies in your application under test) is to run your tests in a random order. Within RSpec this

is an easy thing to accomplish. It's just another command-line argument to pass in at runtime.

```
parallel_rspec --test-options '-r ./config_cloud.rb --order random'
```

This coupled with parallelization will make your tests really work for you.

Chapter 15

Simple Command-line Test Execution

"Am I the only Ruby dev who writes rake tasks for everything? Don't want to worry about the syntax of 50 different cmds or which dir I'm in." [\[source\]](#)

-- [Russ Olsen](#) (author of [Eloquent Ruby](#) and [Design Patterns in Ruby](#))

There's a lot to remember every time we want to run our tests now. So let's clean things up with some Rake tasks.

[Rake](#) is a library in Ruby that enables us to easily capture and run repetitive tasks. To set it up, we just need to install the gem, add a `Rakefile` in our root directory, and add our tasks to the `Rakefile`. Then all we have to do is run our tasks using `rake` from the command-line.

Let's step through an example.

An Example

Part 1: Setup

First we'll want to add Rake to our `Gemfile` and install it (with `bundle install`).

```
# filename: Gemfile
source 'https://rubygems.org'

gem 'rspec', '~> 3.4.0'
gem 'selenium-webdriver', '2.48.1'
gem 'sauce_whisk', '~> 0.0.19'
gem 'parallel_tests', '~> 2.0.0'
gem 'rake', '~> 10.4.2'
```

Next, let's create a `Rakefile` with some initial tasks.


```

# filename: Rakefile
def launch_with(config_filename)
  system("parallel_rspec --test-options '-r ./config/#{config_filename}' --order random"
spec")
end

desc 'Run tests locally'
task :local, :browser do |t, args|
  ENV['browser'] = args[:browser]
  launch_with('local.rb')
end

desc "Run tests in Sauce Labs"
task :cloud do
  launch_with('cloud.rb')
end

```

Notice that this looks similar to a standard Ruby file. We are declaring a method, do some string interpolation in it, referencing that method, and setting some environment variable values. It's the task declaration that's a little bit different (and specific to Rake).

To declare a task, we first want to give it a description. Descriptions are useful since they show up in our task list (next to the command to launch the task). If one is not set, then the task will not appear in the task list. To set one, we use the keyword `desc` and provide a string value (e.g., `'Run tests locally'`).

Once we have that, we can create the task itself. This is done with the keyword `task` followed by the name of the task written as a symbol (e.g., `task :cloud`). The first line of a task ends with the word `do`. This opens up the code block and terminates with the word `end`. Between `do` and `end` is where we place the code for our task.

Our `:local` task is different than the `:cloud` task since we are taking an argument. To accomplish this we specify an additional parameter on the `task` line (e.g., `task :local, :browser`) and some variables in pipes after `do` (e.g., `do |t, args|`). With the `args` variable we are able to receive an argument at run time and pass it into the browser environment variable (e.g., `ENV['browser'] = args[:browser]`).

NOTE: For more information on Rake, check out [this tutorial](#).

In our `launch_in_parallel` method we pull out the lengthy command used to run our tests and wrap it in a `system` command to launch it. We also make this method accept an argument for the config file and inject it into the execution string with interpolation. This enables us to use it for each of our rake tasks. It also keeps our tasks clean, removing redundancy, and enables us to easily change our execution string (if we need to add additional arguments).

If we save this file and run `rake -T` from the command-line, we should see a list of our available Rake actions.

```
rake cloud          # Run tests in Sauce Labs
rake local[browser] # Run tests locally
```

To run a task, just type `rake` followed by the task name. For tasks that take an argument, wrap the argument in square-brackets (e.g., `rake local[chrome]`).

If we run either the `cloud` or `local` tasks, they will run our tests in parallel with 2 processes, in a random order, and will pass. But the cloud task will only run the values specified in `config/cloud.rb` or specified in environment variables at run-time. Let's improve our cloud task.

Part 2: Tasks For The Browsers We Care About

Now that we have Rake tasks we are in a better position to run tests quickly. But in order to make this approach worthwhile, we want to add more tasks to support different browser combinations when running our tests in Sauce Labs.

We can do that pretty easily by replacing our current `:cloud` task with something new. We can specify the browsers we care about in a collection and generate a set of tasks based on that collection. And to make things succinct, we can wrap these tasks in a cloud category (or `namespace`) so they're all helpfully named.

```
# filename: Rakefile
# ...
namespace :cloud do
  BROWSERS_CLOUD = [ 'firefox', 'chrome', 'internet_explorer' ]
  BROWSERS_CLOUD.each do |browser|
    desc "Run tests in Sauce Labs on #{browser.capitalize}"
    task browser.to_sym, :version, :os do |t, args|
      ENV['browser_version'] = args[:version]
      ENV['operating_system'] = args[:os]
      launch_with('cloud.rb')
    end
  end
end
```

In Rake we can group common tasks together with a `namespace` , which follows the same naming conventions as a task. This helps organize our task list, making things more intuitive to run. You'll see what I mean soon.

After specifying each browser we care about as strings in an array, we iterate through the array and create a task based on the browser name (converting it from a string to a symbol). The task accepts two arguments for the browser version and the operating system and passes these values

into the appropriate environment variable.

If we run `rake -T` it will display a list of available tasks.

```
rake cloud:chrome[version,os]           # Run tests in Sauce Labs on Chrome
rake cloud:firefox[version,os]          # Run tests in Sauce Labs on Firefox
rake cloud:internet_explorer[version,os] # Run tests in Sauce Labs on
Internet Explorer
rake local[browser]                     # Run tests locally
```

To run our tests against Internet Explorer 8, all we have to do now is run `rake cloud:internet_explorer['8','Windows XP']` (notice that there are no spaces around the comma). And for different versions of Internet Explorer, we can easily do it by specifying the various available combinations.

```
rake cloud:internet_explorer['6','Windows XP']
rake cloud:internet_explorer['7','Windows XP']
rake cloud:internet_explorer['8','Windows XP']
rake cloud:internet_explorer['8','Windows 7']
rake cloud:internet_explorer['9','Windows 7']
rake cloud:internet_explorer['10','Windows 7']
rake cloud:internet_explorer['10','Windows 8']
```

NOTE: If we try for a browser version and operating system combination that's not valid in Sauce Labs, we'll receive an error message from Sauce Labs stating 'Unsupported OS/browser/version combo'.

Part 3: A Small Bit of Clean-up

Let's apply the same configuration for local execution (replacing our `:local` task).

```
# filename: Rakefile
# ...
namespace :local do
  BROWSERS_LOCAL = [ 'firefox', 'chrome' ]
  BROWSERS_LOCAL.each do |browser|
    desc "Run tests locally on #{browser.capitalize}"
    task browser.to_sym do
      ENV['browser'] = browser
      launch_with('local.rb')
    end
  end
end
# ...
```

Now our task list should look like this.

```
rake cloud:chrome[version,os]           # Run tests in Sauce Labs on Chrome
rake cloud:firefox[version,os]          # Run tests in Sauce Labs on Firefox
rake cloud:internet_explorer[version,os] # Run tests in Sauce Labs on
Internet_explorer
rake local:chrome                       # Run tests locally on Chrome
rake local:firefox                     # Run tests locally on Firefox
```

There are over 650 browser/OS combinations to choose from in Sauce Labs (you can see a full list [here](#)), and there are numerous ways to organize your Rake tasks. So if you aren't looking forward to the prospect of remembering and typing out the operating system, browser, or browser version for your test runs, simply create tasks for each of the browser/OS combinations you care about.

That's one of the biggest advantages to having a Rakefile. You can be as explicit as you want, have a central list to see what is available, and can easily run each item on that list.

Chapter 16

Automating Your Test Runs

You'll probably get a lot of mileage out of your test suite in its current form if you just run things from your computer, look at the results, and inform people when there are issues. But that only helps you solve part of the problem.

The real goal in all of this is to find issues reliably, quickly, and automatically. We've built things to be reliable and quick. Now we need to make them run automatically -- and ideally in sync with the development workflow you are a part of. In order to do that, we want to use a Continuous Integration (CI) server.

A Continuous Integration Primer

Continuous Integration (a.k.a. CI) is the practice of merging code that is actively being worked on into a shared mainline (e.g., trunk or master) as often as possible (e.g., several times a day). This is with the hopes of finding issues early and avoiding merging and integration issues that are not only considered a special kind of hell, but can dramatically slow the time it takes to release software.

The use of a CI server (a.k.a. build server) enables this practice to be automated, and to have tests run as part of the work flow. The lion's share of tests that are typically run on a CI Server are unit (and potentially integration) tests. But we can very easily add in our Selenium tests.

There are numerous CI Servers available for use today, most notably:

- [Bamboo](#)
- [Jenkins](#)
- [Solano Labs](#)
- [TravisCI](#)

Part 1: Tagging & Workflow

In order to get the most out of our test runs in a CI environment we want to break up our test suite into small, relevant chunks and have separate jobs for each. This helps keep test runs fast and informative (so people on your team will care about them). [Gojko Adzic refers to these as "Test Packs"](#).

The workflow is pretty straightforward. The CI Server pulls in the latest code, merges it, and runs unit tests. We then have the CI Server kick off a new job to deploy to a test server and run a

subset of critical acceptance tests (e.g., smoke or sanity tests). Assuming those pass, we can have another job run the remaining tests after that (e.g., the less critical and longer running tests). [Adam Goucher refers to this strategy as a 'shallow' and 'deep' tagging model.](#)

To demonstrate this, let's tag our tests and update our rake tasks to support tags.

An Example

RSpec comes built in with tagging support. It's a simple matter of adding a key/value pair or a symbol to denote what you want. You can place it on individual tests, or a group of tests. And you can use as many tags as you want (separating them with commas).

Let's add some to our specs, following Adam Goucher's shallow and deep approach.

```
# filename: spec/dynamic_loading_spec.rb
# ...
describe 'Dynamic Loading', :deep do
# ...
```

```
# filename: spec/login_spec.rb
# ...
describe 'Login', :shallow do
# ...
```

If we wanted to apply this tag directly to a test, then it would look like this:

```
it 'succeeded', :shallow do
```

To run tests based on a specific tag, we will need to pass in an additional argument to RSpec. It starts with `--tag` followed by the tag value (e.g., `--tag shallow` or `--tag deep`).

Let's update our 'Rakefile' tasks to handle this.

```
# filename: Rakefile
def launch_with(config_filename)
  if ENV['tag']
    test_options = "-r ./config/#{config_filename} --order random --tag #{ENV['tag']}"
  else
    test_options = "-r ./config/#{config_filename} --order random"
  end
  system("parallel_rspec --test-options '#{test_options}' spec")
end
# ...
```

By updating our `launch_in_parallel` method to use a conditional based on the existence of the tag environment variable (`ENV['tag']`), we are able to dynamically alter our execution string at run time without having to change any of our tasks.

We just have to specify the tags at run time before our rake task (similar to how we would specify a different `base_url`).

```
tag=shallow rake local:firefox
```

Part 2: Reporting

In order to make our test output useful for a CI Server we need to generate it in a standard way. One format that works across all CI Servers is JUnit XML.

An Example

This functionality doesn't come built into RSpec, but it's simple enough to add through the use of a third-party library. There are plenty to choose from, but we'll go with [rspec_junit_formatter](#).

```
# filename: Gemfile
source 'https://rubygems.org'

gem 'rspec', '~> 3.4.0'
gem 'selenium-webdriver', '2.48.1'
gem 'sauce_whisk', '~> 0.0.19'
gem 'parallel_tests', '~> 2.0.0'
gem 'rake', '~> 10.4.2'
gem 'rspec_junit_formatter', '~> 0.2.3'
```

After we install it we need to specify the formatter type and an output file for RSpec to consume. This is done through the use of two new arguments; `--format RspecJunitFormatter` and `--out results.xml` . But in order to use this with our parallel test execution we need to create a uniquely named XML output file for each test process. To accomplish that with `parallel_tests` we will need to take advantage of another feature in RSpec -- the `.rspec` file.

RSpec comes with the ability to place command arguments in a `.rspec` file within the root directory of our project, which will automatically be consumed at runtime. So we'll place our new arguments there, and use a kind of interpolation to inject an environment variable from `parallel_tests` which will be unique for each test process.

We're doing this instead of adding it to the `launch_in_parallel` method in our 'Rakefile' because of how `parallel_tests` handles execution. In order to get dynamically named XML output, we have

to use the `.rspec` file in order to get the timing right. Also, we don't want our XML output to happen all the time. So let's also make it so the commands in `.rspec` only run when we want them to.

```
# filename: .rspec
<% if ENV['ci'] == 'on' %>
--format RspecJunitFormatter
--out results/result<%= ENV['TEST_ENV_NUMBER'] %>.xml
<% end %>
```

When we run our tests with the `ci` environment variable set to `'on'` (e.g., `ci=on rake local:firefox`), each test process creates its own `result.xml` by number (e.g., `result.xml`, `result2.xml`, etc.), and ends up in a `results` directory.

After a test run our directory structure will look like this:


```
.
Gemfile
Gemfile.lock
Rakefile
config
  cloud.rb
  local.rb
pages
  base_page.rb
  dynamic_loading.rb
  login.rb
results
  result.xml
  result2.xml
spec
  dynamic_loading_spec.rb
  login_spec.rb
  spec_helper.rb
vendor
  chromedriver
  sauce-connect
    bin
      sc
      sc.dSYM
        Contents
          Info.plist
          Resources
            DWARF
              sc
    include
      sauceconnect.h
    lib
      libsauceconnect.a
      libsauceconnect.la
    license.html
```

If we open one of the XML files from the `results` directory, it will contain a bunch of info from the test run. This is what our CI server will use to track test results (e.g., passes, failures, time to complete, etc.) and trend them over time.

Here's what one from a passing run looks like.

```
<?xml version="1.0" encoding="UTF-8"?>
<testsuite name="rspec" tests="2" failures="0" errors="0" time="25.890860" timestamp=
"2015-11-21T11:35:05-05:00">
  <!-- Randomized with seed 5604 -->
  <properties/>
  <testcase classname="spec.dynamic_loading_spec" name="Dynamic Loading Example 1:
Hidden Element" file="./spec/dynamic_loading_spec.rb" time="13.115928"/>
  <testcase classname="spec.dynamic_loading_spec" name="Dynamic Loading Example 2:
Rendered after the fact" file="./spec/dynamic_loading_spec.rb" time="12.774309"/>
</testsuite>
```

We don't want to commit the XML files our repository, so let's add them to our ignore file.

```
# filename: .gitignore

*.xml
```

In this case we're using [Git](#). If you're using something else (e.g., [Mercurial](#), [Subversion](#), etc.) then find its ignore file equivalent and use it.

Now we're ready to wire things up to our CI server.

Part 3: CI Server Job Configuration

[Jenkins](#) is a fully functional, widely adopted, and open-source CI server. It's a great candidate for us to step through.

Let's start by setting it up on the same machine as our test code. Keep in mind that this isn't the "proper" way to go about this — it's merely beneficial for this example. To do it right, the Jenkins server (e.g., master node) would live on a machine of its own.


1. Quick Setup

A simple way to get started is to grab the latest Jenkins war file. You can grab it from the [Jenkins homepage](#), or from [the direct download link on the homepage](#).

Once downloaded, launch it from the terminal.

```
> java -jar jenkins.war
...
INFO: Jenkins is fully up and running
```

You will now be able to use Jenkins by visiting <http://localhost:8080/> in your browser.



NOTE: Before moving to the next step, click **ENABLE AUTO-REFRESH** at the top right-hand side of the page. Otherwise you'll need to manually refresh the page (e.g., when running a job and waiting for results to appear).

2. Job Creation and Configuration

Now that Jenkins is loaded in the browser, let's create a Job and configure it to run our Shallow tests against an old version of Internet Explorer (e.g., IE8).

- Click **New Item** from the top-left of the Dashboard
- Give it a name (e.g., **Shallow Tests IE8**)
- Select the **Freestyle Project** option
- Click **OK**

New Item [Jenkins] x

localhost:8080/newJob

Jenkins

search

Jenkins

- New Item
- People
- Build History
- Manage Jenkins
- Credentials

Build Queue

No builds in the queue.

Build Executor Status

1	Idle
2	Idle

Item name

☒ **Freestyle project**
This is the central feature of Jenkins. Jenkins will build your project, combining any SCM with any build system, and this can be even used for something other than software build.

☐ **Maven project**
Build a maven project. Jenkins takes advantage of your POM files and drastically reduces the configuration.

☐ **External Job**
This type of job allows you to record the execution of a process run outside Jenkins, even on a remote machine. This is designed so that you can use Jenkins as a dashboard of your existing automation system. See [the documentation for more details](#).

☐ **Multi-configuration project**
Suitable for projects that need a large number of different configurations, such as testing on multiple environments, platform-specific builds, etc.

OK

[Help us localize this page](#) Page generated: Nov 28, 2015 1:30:07 PM [REST API](#) [Jenkins ver. 1.638](#)

This will load a configuration screen for the Jenkins job.

The screenshot shows the Jenkins web interface in a browser window. The address bar shows 'localhost:8080/job/Shallow%20Tests%20IE8/configure'. The Jenkins logo and a search bar are at the top. A breadcrumb trail shows 'Jenkins > Shallow Tests IE8 > configuration'. On the left sidebar, there are links: 'Back to Dashboard', 'Status', 'Changes', 'Workspace', 'Build Now', 'Delete Project', and 'Configure'. Below these is a 'Build History' section with a search box and RSS links. The main configuration area has a 'Project name' field with 'Shallow Tests IE8' and a 'Description' text area. Below the description is a '[Plain text] Preview' link. There are four checkboxes: 'Discard Old Builds', 'This build is parameterized', 'Disable Build (No new builds will be executed until the project is re-enabled.)', and 'Execute concurrent builds if necessary'. Each checkbox has a help icon. Below these is the 'Advanced Project Options' section with an 'Advanced...' button. The 'Source Code Management' section has radio buttons for 'None' (selected), 'CVS', 'CVS Projectset', and 'Subversion'. The 'Build Triggers' section has checkboxes for 'Build after other projects are built', 'Build periodically', and 'Poll SCM', each with a help icon. The 'Build' section has an 'Add build step' button. The 'Post-build Actions' section has 'Save' and 'Apply' buttons.

Shallow Tests IE8 Config

localhost:8080/job/Shallow%20Tests%20IE8/configure

Jenkins

search

Jenkins > Shallow Tests IE8 > configuration

[Back to Dashboard](#)

[Status](#)

[Changes](#)

[Workspace](#)

[Build Now](#)

[Delete Project](#)

[Configure](#)

Build History [trend](#)

find x

[RSS for all](#) [RSS for failures](#)

Project name: Shallow Tests IE8

Description: [Text area]

[Plain text] [Preview](#)

☐ Discard Old Builds ?

☐ This build is parameterized ?

☐ Disable Build (No new builds will be executed until the project is re-enabled.) ?

☐ Execute concurrent builds if necessary ?

Advanced Project Options [Advanced...](#)

Source Code Management

☒ None

☐ CVS

☐ CVS Projectset

☐ Subversion

Build Triggers

☐ Build after other projects are built ?

☐ Build periodically ?

☐ Poll SCM ?

Build

[Add build step](#)

Post-build Actions

[Save](#) [Apply](#)

- Under **Advanced Project Options** click **Advanced...**
- Check the box next to **Use custom workspace**
- In the input field next to **Directory** input the full path to your test code

NOTE: Ideally, your test code would live in a version control system and you would configure your job (under **source code Management**) to pull it in and run it. To use this approach you may need to install a plugin to handle it (depending on your version control solution). For more info on plugins in Jenkins, go [here](#).

The screenshot shows the Jenkins web interface in a browser. The page title is 'Shallow Tests IE8 Config'. The URL is 'localhost:8080/job/Shallow%20Tests%20IE8/configure'. The Jenkins logo and a search bar are at the top. The left sidebar contains links: 'Back to Dashboard', 'Status', 'Changes', 'Workspace', 'Build Now', 'Delete Project', and 'Configure'. Below these is a 'Build History' section with a search bar and RSS links. The main content area is the configuration page for 'Shallow Tests IE8'. It has a 'Project name' field with the value 'Shallow Tests IE8' and a 'Description' text area. Below these are several checkboxes: 'Discard Old Builds', 'This build is parameterized', 'Disable Build (No new builds will be executed until the project is re-enabled.)', and 'Execute concurrent builds if necessary'. There is an 'Advanced Project Options' section with checkboxes for 'Quiet period', 'Retry Count', 'Block build when upstream project is building', 'Block build when downstream project is building', and 'Use custom workspace' (which is checked). The 'Use custom workspace' section has a 'Directory' field with the path '/Users/tourdedave/Dropbox/_dev/selenium-guidebook/code-examples-ruby/16_ci/' and a 'Display Name' field. There is also a 'Source Code Management' section with radio buttons for 'None' (selected), 'CVS', and 'CVS Projectset'. At the bottom are 'Save' and 'Apply' buttons.

- Scroll down until you reach the **Build** section (near the bottom of the page)
- Click on **Add build step** and select **Execute Shell**

Build

The screenshot shows the 'Build' section of the Jenkins configuration page. At the top is a button labeled 'Add build step' with a dropdown arrow. A dropdown menu is open, showing four options: 'Execute Windows batch command', 'Execute shell' (which is highlighted in blue), 'Invoke Ant', and 'Invoke top-level Maven targets'. Below the dropdown menu are 'Save' and 'Apply' buttons.

In the Command input box, add the following commands:

```
export SAUCE_USERNAME="your-sauce-username"
export SAUCE_ACCESS_KEY="your-sauce-access-key"
gem install bundler
bundle install
ci=on bundle exec rake cloud:internet_explorer['8','Windows XP']
```

Build

Execute shell

Command

```
cd /Users/tourdedave/Dropbox/_dev/selenium-guidebook/code-examples-ruby/16_ci/02
export SAUCE_USERNAME="your-sauce-username"
export SAUCE_ACCESS_KEY="your-sauce-access-key"
gem install bundler
bundle install
ci=on bundle exec rake cloud:internet_explorer['8','Windows XP']
```

See [the list of available environment variables](#)

Delete

Since our tests have never run on this server we need to include the installation and running of the bundler gem (`gem install bundler` and `bundle install`) to download and install the libraries (a.k.a. gems) used in our test suite. And we also need to specify our credentials for Sauce Labs (unless you decided to hard-code these values in `config/cloud.rb` - if so, then you don't need to specify them here).

Now we need to consume the JUnit XML that our test suite will generate.

- Click on `Add post-build action`
- Select `Publish JUnit test result report`

Build

Execute shell

Command

```
gem install bundler
bundle install
ci=on bundle exec rake cloud:internet_explorer['8','Windows XP']
```

- Aggregate downstream test results
- Archive the artifacts
- Build other projects
- Publish JUnit test result report**
- Publish Javadoc
- Record fingerprints of files to track usage
- E-mail Notification

Add post-build action ▼

Save

Apply

- In the input box to the right of `Test Report XMLs` type `results/*.xml`

Post-build Actions

Publish JUnit test result report

Test report XMLs

[Fileset 'includes'](#) setting that specifies the generated raw XML report files, such as 'myproject/target/test-reports/*.xml'. Basedir of the fileset is [the workspace root](#).

☐ Retain long standard output/error 

Health report amplification factor

1% failing tests scores as 99% health. 5% failing tests scores as 95% health

Delete

Now we're ready to save, run our tests, and view the job result.

- Click **Save**
- Click **Build Now** from the left-hand side of the screen

When the build completes, the result will be listed on the job's home screen. In this case, the job passed.



NOTE: If you had a different result, you can drill into a job to see what was happening behind the scenes. To do that click on the build you want from **Build History** and select **Console Output**.

This output will be your best bet in tracking down an unexpected result.

A passing job means passing tests, which is great. But we'll also want to see what a failure looks like to make sure its helpful.


Part 3: Force A Failure

Let's add a new test to `spec/login_spec.rb` that will fail every time we run it.

```
# filename: spec/login_spec.rb
# ...
it 'forced failure' do
  expect(false).to eq true
end

end
```

Now when we run our Jenkins job again, it will fail.



The screenshot shows the Jenkins web interface in a browser. The address bar indicates the URL is `localhost:8080/job/Shallow%20Tests%20IE8/8/`. The Jenkins logo is visible at the top left. The main content area displays the details for **Build #2 (Nov 28, 2015 2:34:44 PM)**. On the left sidebar, there are links for [Back to Project](#), [Status](#), [Changes](#), [Console Output](#), [Edit Build Information](#), [Delete Build](#), [Test Result](#), and [Previous Build](#). The main area shows a yellow circle icon for the build, a notepad icon with the text "No changes.", and an orange diamond icon with the text "Started by anonymous user". Below these, there is a link for [Test Result \(1 failure / +1\)](#) with the subtext [spec.login_spec.Login forced failure](#). On the right side, it says "Started 16 min ago" and "Took 46 sec", with a link for [add description](#). At the bottom of the page, there is a footer with a link for [Help us localize this page](#), the text "Page generated: Nov 28, 2015 2:51:04 PM", and links for [REST API](#) and [Jenkins ver. 1.638](#).

When we click on the failed test we can see the failure message along with a URL to the job in Sauce Labs.

Shallow Tests IE8 #8 test - x

localhost:8080/job/Shallow%20Tests%20IE8/8/testReport/junit/spec/login_spec/Login_forced_failure/

Jenkins


search

Jenkins > Shallow Tests IE8 > #2 > Test Results > Test Results > spec > login_spec > Login forced failure [DISABLE AUTO REFRESH](#)

- Back to Project
- Status
- Changes
- Console Output
- Edit Build Information
- History
- Test Result
- Previous Build

Failed

spec.login_spec.Login forced failure (from rspec2)

Failing for the past 1 build (Since  #8)
Took 7.7 sec.
[add description](#)

Error Message

RSpec::Core::MultipleExceptionError

Stacktrace


```
expected: true
got: false

(compared using eql?)
```

Watch a video of the test at <https://saucelabs.com/tests/d5e2048cff204a8fbb52e127e64b6617>

[Help us localize this page](#) Page generated: Nov 28, 2015 2:41:19 PM [REST API](#) [Jenkins ver. 1.638](#)

When we follow the URL to the Sauce Labs job we're able to see what happened during the test run (e.g., we can replay a video of the test, see what Selenium commands were issued, etc.).



Login forced failure

failed

by the-internet

Start Time: 11-28-2015 14:34:50

| Duration: 9s

Visibility: public

Report Issue

XP 8 Not on Sauce Connect

All tests

Commands	Screenshot	Selenium Log	Metadata
POST /session desiredCapabilities: {"rotatable":false,"takesScreenshot":true,"nativeEvent... => {"takesScreenshot":true,"unexpectedAlertBeha...	0 (+1.58s)		
POST url url: "http://the-internet.herokuapp.com/login" => null	4.28s (+2.13s)		
POST element using: "id" value: "login" => {"ELEMENT":"0"}	4.72s (+0.09s)		
GET element/0/displayed => true	4.98s (+0.08s)		
DELETE /session/d5e2048cff204a8fbb52e1 07-546-6617	5.64s (+0s)		

SCREENCAST

Notifications

In order to maximize your CI effectiveness, you'll want to send out notifications to alert your team members when there's a failure.

There are numerous ways to go about this (e.g., e-mail, chat, text, co-located visual cues, etc). And thankfully there are numerous, freely available plugins that can help facilitate whichever method you want. You can find out more about Jenkins plugins [here](#).

For instance, if you wanted to use chat notifications and you use a service like HipChat or Slack, you would do a plugin search and find the following plugins:

Filter:

Updates Available Installed Advanced

Install ↓	Name	Version
<input checked="" type="checkbox"/>	HipChat Plugin This plugin allows your team to setup build notifications to be sent to HipChat rooms.	0.1.8

Install without restart Download now and install after restart

Filter:

Updates Available Installed Advanced

Install ↓	Name	Version
<input type="checkbox"/>	Slack Notification Plugin A Build status publisher that notifies channels on a Slack team	1.7

Install without restart Download now and install after restart

After installing the plugin for your chat service, you will need to provide the necessary information to configure it (e.g., an authorization token, the channel/chat room where you want notifications to go, what kinds of notifications you want sent, etc.) and then add it as a `Post-build Action` to your job (or jobs).

After installing and configuring a plugin, when your CI job runs and fails, a notification will be sent to the chat room you configured.

Outro

By using a CI Server you're able to put your tests to work by using computers for what they're good at, automation. This frees you up to focus on more important things. But keep in mind that there are numerous ways to configure your CI server. Be sure to tune it to what works best for you and your team. It's well worth the effort.

Chapter 17

Finding Information On Your Own

There is information all around us when it comes to Selenium. But it can be challenging to sift through it, or know where to look.

Here is a list breaking down a majority of the Selenium resources available, and what they're useful for.

Documentation & Tips

- [Selenium HQ](#)

This is the official Selenium project documentation site. It's a bit dated, but there is loads of helpful information here. You just have to get the hang of how to navigate the site to find what you need.

- [The Selenium Wiki](#)

This is where all the good stuff is -- mainly, documentation about the various language bindings and browser drivers. If you're not already familiar with it, take a look.

- [Elemental Selenium Archives](#)

Every tip I've written is freely available on the tips archive page. There are over 70 different Selenium problems and solutions covered.

Blogs

- [The official Selenium blog](#)

This is where news of the Selenium project gets announced, and there's also the occasional round-up of what's going on in the tech space (as it relates to testing). Definitely worth a look.

- [A list of "all" Selenium WebDriver blogs](#)

At some point, someone rounded up a large list of blogs from Selenium practitioners and committers. It's a pretty good list.

Other Books

- [Selenium 2 Testing Tools](#)

This book is by [David Burns](#) and it is a solid resource. It outlines how to use Selenium, top-to-bottom. It's definitely worth having on your shelf.

- [Selenium Testing Tools Cookbook](#)

This is another good book outlining some great ways to leverage Selenium. While I haven't had a chance to finish reading it, it's clear that Gundecha has a very pragmatic approach that will yield great results.

- [Selenium Design Patterns and Best Practices](#)

Dima Kovalenko's book covers useful tactics and strategies for successful test automation with Selenium. I was a technical reviewer for the book and think it's a tremendous resource. The book covers Ruby, and he has ported the examples to Java. You can find them [here](#).

Meetups

- [The Selenium Hangout](#)

This is an entirely online meetup that is run by me and [David Burns](#) where we talk about the latest in the Selenium community (e.g., upcoming conferences, the status of Selenium 3, the W3C spec, etc.) and answer questions from the community. Videos are recorded and made available (along with notes) on [the official Selenium blog](#).

- [All Selenium Meetups listed on Meetup.com](#)

A listing of all in-person Selenium Meetups are available on Meetup.com. If you're near a major city, odds are there's one waiting for you.

- [How to start your own Selenium Meetup](#)

If there's not a Selenium Meetup near you, start one! Sauce Labs has a great write up on how to do it.

Conferences

- [Selenium Conf](#)

This is the official annual conference of the Selenium project where practitioners and committers gather and share their latest knowledge and experiences with testing. The conference location changes every year (e.g., it's been in San Francisco, London, Boston, Bangalore, and soon -- Portland).

- [Selenium Camp](#)

This is an annual Selenium conference in Eastern Europe (in Kiev, Ukraine) organized by the folks

at [XP Injection](#). It's a terrific conference. If you can make the trip, I highly recommend it.

- [List of other testing conferences](#)

[Alister Scott](#) has a good write-up on the testing conferences he thought about attending last year.

Videos

- [Selenium Conference Talks](#)

All of the talks from The Selenium Conference are recorded and made freely available online. This is a tremendous resource.

- [Selenium Meetup Talks](#)

Some of the Selenium Meetups make it a point to record their talks and publish them afterwards. Here are some of them. They are a great way to see what other people are doing and pick up some new tips.

- [Selenium Hangout](#)

All of the Selenium Hangout Meetups are recorded and made available. A lot of great stuff is discussed in them.

Mailing Lists

- [Selenium Developers List](#)

This is where developers discuss changes to the Selenium project.

- [Selenium Users Google Group](#)
- [Selenium LinkedIn Users Group](#) The signal to noise ratio in these groups can be challenging at times. But you can occasionally find some answers to your questions.
- [Agile Testing Yahoo Group](#)

This list is less about Selenium and more about "testing" (in the [checking versus testing](#) sense). But it's a great place to ask for sage advice, and to hang out and listen. You will learn some valuable insights here.

Forums

- [Stack Overflow](#)
- [Quora](#)

These are the usual forums where you can go looking for answers to questions you're facing (in addition to the mailing lists above).

Issues

- [Selenium Issue Tracker](#)

If you're running into a specific and repeatable issue that just doesn't make sense, you may have found a bug in Selenium. You'll want to check the Selenium Issue Tracker to see if it has already been reported. If not, then create a new issue. But be sure to read [this post](#) before you do (so you can be sure that you have provided enough information).

Chatting With the Selenium Community on IRC

The Selenium IRC Chat Channel is arguably the best way to connect with the Selenium community and get questions answered. This is where committers and practitioners hang out day-in and day-out.

Brief Intro To IRC

IRC (short for Internet Relay Chat) is a protocol that freely enables live chatting (both in groups and person to person) and file sharing. It's been around for a while (circa 1988) and is the preferred mode of communication among certain tech circles.

Within the realm of IRC there are numerous networks you can connect to. Each one containing people and bots logged in and joined to one or more chat rooms talking, sharing files, etc.

One of the beautiful things about IRC is that there is no registration required to join the party. You just need to download a client that handles the IRC protocol (there's at least one available for every platform), point it at a network, and pick a unique nickname for yourself on that network.

Once you're on you can join a chat room and start jib-jabbing.

How To Get Connected

Step 1: Get An IRC Chat Client

First thing's first, get a chat client that supports IRC.

You may already have one and not even know. For example, [Adium](#) (for OSX) supports a staggering number of chat protocols. If you already have it (or something like it) then use that to connect. If you don't, then you'll need to download one that supports it (or one that is built specifically for IRC).

Here are some worthwhile IRC chat clients (broken out by platform -- and are free unless otherwise noted).

OSX

- [LimeChat](#)
- [Textual](#) (\$4.99 to buy)

Windows

- [mIRC](#) (free for 30 days, \$20 for a single-user license)

Linux

- [Irssi](#)

Web-only

- [Webchat](#)
- [IRCCloud](#) (free with paid tier as well, Android and iOS apps available as well)

Step 2: Connect To The Proper Server

The Selenium chat channel lives on the Freenode network. To connect to it directly you would use `irc.freenode.net`.

Before connecting you should be able to set a nickname and perhaps even specify which channel you would like to connect to after connecting (e.g., `#selenium`). If you don't see these bits, don't sweat it. Connect and proceed to the next step.

NOTE: If you're using Webchat, it will automatically connect you to Freenode.

Step 3: Join The Chat Channel

In IRC parlance channels are prepended with a `#` and are lower-case. So the Selenium channel is `#selenium`.

If you were able to configure your chat client to join the channel for you, then proceed to the next step. Otherwise, you'll need to issue a command in the status window. There are a series of commands you can issue in IRC. They are all prepended with a `/`. To join a chat channel type `/join #selenium` and press `Enter`.

This will open a new chat window to the Selenium chat channel. Woohoo!

Step 4: Talk And Hangout

Feel free to say hello and introduce yourself. But more importantly, ask your question. If it looks like no one is chatting, ask it anyway. Someone will see it and eventually respond. They always do.

In order to get your answer, you'll probably need to hang around for a bit. But the benefit of

being a fly on the wall is that you gain insight into other problems people face, possible solutions, and the current state of the Selenium project and its various pieces.

Chapter 18

Now You Are Ready

The journey for doing Selenium successfully can be long and arduous. But by adhering to the principals in this book, you will avoid a majority of the pitfalls around you. You're also in a better position now -- armed with all of the information necessary to continue your Selenium journey.

You are ready. Keep going, and best of luck!

And if you have any questions, feedback, or want help -- get in touch!

- E-mail: dhaeffner@gmail.com
- Twitter: [@TourDeDave](https://twitter.com/TourDeDave)
- Office hours: [SoHelpful.me/TourDeDave](https://sohelpful.me/TourDeDave)

My office hours are free (read: completely free -- no strings attached). You can grab a 30-minute time slot to hop on a Skype or Google Hangout call with me to talk about testing. And if we're both at a conference or meetup, feel free to stop me and say hello!